

Exercise

10

TUM Department of Informatics

Supervised by	Prof. Dr. Stephan Günnemann Informatics 3 - Professorship of Data Mining and Analytics
Submitted by	Marcel Bruckner (03674122) Julian Hohenadel (03673879) Kevin Bein (03707775)
Submission date	Munich, January 2, 2020

Deep Learning II

Problem 1:

a) Since all weights are positive, $\max(0, W_j^{*T} x_i) = W_j^{*T} x_i$ and the following holds:

$$\begin{aligned}\mathcal{L}_{NN}(W_{NN}^*) &= \frac{1}{2} \sum_{i=1}^N (-y_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (W_{L+1}^{*T} (W_L^{*T} \max(0, \dots (0, W_2^{*T} \max(0, W_1^{*T} x_i)))) - y_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (W_{L+1}^{*T} (W_L^{*T} (\dots (W_2^{*T} (W_1^{*T} x_i)))) - y_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (W_{L+1}^{*T} \cdot \dots \cdot W_1^{*T} x_i - y_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^N (W_{NN}^{*T} x_i - y_i)^2\end{aligned}$$

Since the solutions W_{NN}^* and w_{LS}^* are a global optimum, we can set $W_{NN}^* = w_{LS}^*$ and we get $\mathcal{L}_{NN}(W_{NN}^*) = \mathcal{L}_{LS}(w_{LS}^*)$.

b) Since simple linear regression is a special case of a Feed-Forward Neural Network, the Loss is the same. This does not work the other way around because NNs can learn more complex functions. w_{LS}^* being non-negative does not imply anything about the optimal weights of the network W_{NN}^* . Therefore we can conclude that $\mathcal{L}_{NN}(W_{NN}^*) \leq \mathcal{L}_{LS}(w_{LS}^*)$.

Problem 2:

exercise_10_notebook_90_36

January 5, 2020

```
[0]: import copy
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
```

```
[2]: !nvidia-smi
```

Fri Jan 3 19:33:22 2020

```
+-----+
| NVIDIA-SMI 440.44          Driver Version: 418.67          CUDA Version: 10.1     |
+-----+-----+-----+
| GPU  Name            Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====+=====+=====+
|   0   Tesla P100-PCIE...    Off   | 00000000:00:04:0 Off  |             0        |
| N/A   35C    P0      28W / 250W |      0MiB / 16280MiB |           0%      Default |
+-----+-----+-----+

+-----+
| Processes:                                     GPU Memory |
|  GPU       PID    Type    Process name                     Usage      |
|=====+=====+
|  No running processes found
+-----+
```

1 PyTorch

In this notebook you will gain some hands-on experience with [PyTorch](#), one of the major frameworks for deep learning. To install PyTorch run `conda install pytorch torchvision cudatoolkit=10.1 -c pytorch`, with cudatoolkit set to whichever CUDA version you have installed. You can check this by running `nvcc --version`. If you do not have an Nvidia GPU you can run `conda install pytorch torchvision cpuonly -c pytorch` instead. However, in this case we recommend using [Google Colab](#).

You will start by re-implementing some common features of deep neural networks (dropout and

batch normalization) and then implement a very popular modern architecture for image classification (ResNet) and improve its training loop.

2 1. Dropout

Dropout is a form of regularization for neural networks. It works by randomly setting activations (values) to 0, each one with equal probability p . The values are then scaled by a factor $\frac{1}{1-p}$ to conserve their mean.

Dropout effectively trains a pseudo-ensemble of models with stochastic gradient descent. During evaluation we want to use the full ensemble and therefore have to turn off dropout. Use `self.training` to check if the model is in training or evaluation mode.

```
[0]: class Dropout(nn.Module):
    """
    Dropout, as discussed in the lecture and described here:
    https://pytorch.org/docs/stable/nn.html#torch.nn.Dropout

    Args:
        p: float, dropout probability
    """
    def __init__(self, p):
        super().__init__()
        self.p = p

    def forward(self, input):
        """
        The module's forward pass.
        This has to be implemented for every PyTorch module.
        PyTorch then automatically generates the backward pass
        by dynamically generating the computational graph during
        execution.

        Args:
            input: PyTorch tensor, arbitrary shape

        Returns:
            PyTorch tensor, same shape as input
        """
        if self.training:
            mask = np.random.random(input.shape)
            return torch.from_numpy(np.where(mask <= self.p, 0, (1 / (1 - self.
↪p))))

        # TODO: Set values randomly to 0.
```

```
[4]: # Test dropout
test = torch.ones(10_000)
dropout = Dropout(0.5)
test_dropped = dropout(test)

print(test_dropped.sum().item())
print((test_dropped > 0).sum().item())

print(np.isclose(test_dropped.sum().item(), 10_000, atol=400))
print(np.isclose((test_dropped > 0).sum().item(), 5_000, atol=200))

# These assertions can in principle fail due to bad luck, but
# if implemented correctly they should almost always succeed.
assert np.isclose(test_dropped.sum().item(), 10_000, atol=400)
assert np.isclose((test_dropped > 0).sum().item(), 5_000, atol=200)
```

```
10104.0
5052
True
True
```

3 2. Batch normalization

Batch normalization is a trick used to smoothen the loss landscape and improve training. It is defined as the function

$$y = \frac{x - \mu_x}{\sigma_x + \epsilon} \cdot \gamma + \beta$$

, where γ and β are learnable parameters and ϵ is some small number to avoid dividing by zero. The statistics μ_x and σ_x are taken separately for each feature. In a CNN this means averaging over the batch and all pixels.

```
[0]: class BatchNorm(nn.Module):
    """
    Batch normalization, as discussed in the lecture and similar to
    https://pytorch.org/docs/stable/nn.html#torch.nn.BatchNorm1d

    Only uses batch statistics (no running mean for evaluation).
    Batch statistics are calculated for a single dimension.
    Gamma is initialized as 1, beta as 0.

    Args:
        num_features: Number of features to calculate batch statistics for.
    """
    def __init__(self, num_features):
        super().__init__()

        # TODO: Initialize the required parameters
```

```

        self.gamma = nn.Parameter(torch.ones(num_features)).unsqueeze(0).
↪unsqueeze(-1)
        #self.gamma = self.gamma.
        self.beta = nn.Parameter(torch.zeros(num_features)).unsqueeze(0).
↪unsqueeze(-1)
        #self.beta = self.beta.

    def forward(self, input):
        """
        Batch normalization over the dimension C of (N, C, L).

        Args:
            input: PyTorch tensor, shape [N, C, L]

        Return:
            PyTorch tensor, same shape as input
        """
        eps = 1e-5

        mean = input.mean(dim=[0, 2], keepdim=True)
        input_mean_norm = input - mean
        var = torch.sqrt(input.var(dim=[0, 2], keepdim=True))

        return (input_mean_norm / (var + eps)) * self.gamma + self.beta

    # TODO: Implement the required transformation

```

```

[6]: # Tests the batch normalization implementation
torch.random.manual_seed(42)
test = torch.randn(8, 2, 4)

b1 = BatchNorm(2)
test_b1 = b1(test)

b2 = nn.BatchNorm1d(2, affine=False, track_running_stats=False)
test_b2 = b2(test)

print(test_b1)
print("-----")
print(test_b2)

print(torch.allclose(test_b1, test_b2, rtol=0.02))
assert torch.allclose(test_b1, test_b2, rtol=0.02)

```

```

tensor([[[ 1.6380,  1.2470,  0.7253, -1.9484],
         [ 0.6971, -1.2133, -0.0234, -1.5829]],

```

```

[[[-0.7447,  1.3905, -0.4249, -1.3242],
  [-0.7073, -0.5391, -0.7482,  0.7810]],

[[ 1.3849, -0.2178, -0.5182,  0.3152],
  [-0.7375,  1.0964,  0.8193,  1.6979]],

[[ 1.0618,  1.0772,  0.4671,  1.1113],
  [-0.2117,  0.0613, -0.2317,  0.8782]],

[[[-1.3073, -0.8507, -0.2745,  1.4516],
  [ 0.3380, -0.4044,  0.3249, -0.7540]],

[[-1.4611,  0.8097, -0.8583, -0.6105],
  [-1.2529,  2.1395, -1.2134, -0.4677]],

[[-0.8886, -0.6611, -0.0064,  0.3918],
  [-0.4678,  1.2093, -0.7933, -0.7154]],

[[-1.3238, -0.0438, -0.1323,  0.5251],
  [-0.0781,  1.8616, -1.1634,  1.4012]]], grad_fn=<AddBackward0>)
-----
tensor([[[ 1.6642,  1.2669,  0.7369, -1.9796],
          [ 0.7082, -1.2327, -0.0238, -1.6083]],

[[[-0.7567,  1.4128, -0.4317, -1.3454],
  [-0.7186, -0.5477, -0.7602,  0.7935]],

[[ 1.4070, -0.2212, -0.5265,  0.3202],
  [-0.7493,  1.1140,  0.8324,  1.7251]],

[[ 1.0788,  1.0945,  0.4746,  1.1291],
  [-0.2151,  0.0622, -0.2354,  0.8923]],

[[[-1.3282, -0.8643, -0.2789,  1.4748],
  [ 0.3434, -0.4109,  0.3301, -0.7660]],

[[-1.4845,  0.8227, -0.8720, -0.6202],
  [-1.2729,  2.1737, -1.2328, -0.4752]],

[[-0.9028, -0.6717, -0.0065,  0.3981],
  [-0.4753,  1.2287, -0.8060, -0.7269]],

[[[-1.3450, -0.0445, -0.1344,  0.5335],
  [-0.0794,  1.8914, -1.1820,  1.4237]]]])
True

```

4 3. ResNet

ResNet is the model that first introduced residual connections (a form of skip connections). It is a rather simple, but successful and very popular architecture. In this part of the exercise we will re-implement it step by step.

Note that there is also an [improved version of ResNet](#) with optimized residual blocks. Here we will implement the [original version](#) for CIFAR-10.

This is just a convenience function to make e.g. `nn.Sequential` more flexible. It is e.g. useful in combination with `x.squeeze()`.

```
[0]: class Lambda(nn.Module):
    def __init__(self, func):
        super().__init__()
        self.func = func

    def forward(self, x):
        return self.func(x)
```

We begin by implementing the residual blocks. The block is illustrated by this sketch:

attachment:residual_connection.png

Note that we use ‘SAME’ padding, no bias, and batch normalization after each convolution.

```
[0]: class ResidualBlock(nn.Module):
    """
    The residual block used by ResNet.

    Args:
        in_channels: The number of channels (feature maps) of the incoming
        ↪ embedding
        out_channels: The number of channels after the first convolution
        stride: Stride size of the first convolution, used for downsampling
    """

    def __init__(self, in_channels, out_channels, stride=1):
        super().__init__()
        if stride > 1 or in_channels != out_channels:
            # Add strides in the skip connection and zeros for the new channels.
            self.skip = Lambda(lambda x: F.pad(x[:, :, ::stride, ::stride],
            ↪ (0, 0, 0, 0, 0, 0, out_channels -
            ↪ in_channels),
            mode="constant", value=0))
        else:
            self.skip = nn.Sequential()

        # TODO: Initialize the required layers
```



```

        #self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
        ↪padding="SAME", bias=False, stride=stride)
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
        ↪padding=1, bias=False, stride=stride, groups=1)
        self.bn1 = nn.BatchNorm2d(out_channels) #BatchNorm(out_channels)
        #self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
        ↪padding="SAME", bias=False, stride=stride)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
        ↪padding=1, bias=False, stride=1, groups=1)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.relu = nn.ReLU()

    def forward(self, input):
        # TODO: Execute the required layers and functions
        #print("input.shape:", input.shape)
        identity = self.skip(input)
        #print("identity.shape:", identity.shape)

        out = self.conv1(input)
        out = self.bn1(out)
        out = self.relu(out)
        #print("out1.shape:", out.shape)

        out = self.conv2(out)
        out = self.bn2(out)

        #print("out2.shape:", out.shape)
        out += identity
        out = self.relu(out)

        #print("out.shape:", out.shape)
        #print("")
        return out

```

Next we implement a stack of residual blocks for convenience. The first layer in the block is the one changing the number of channels and downsampling. You can use `nn.ModuleList` to use a list of child modules.

```

[0]: class ResidualStack(nn.Module):
    """
    A stack of residual blocks.

    Args:
        in_channels: The number of channels (feature maps) of the incoming
        ↪embedding
        out_channels: The number of channels after the first layer
        stride: Stride size of the first layer, used for downsampling

```

```

        num_blocks: Number of residual blocks
    """

    def __init__(self, in_channels, out_channels, stride, num_blocks):
        super().__init__()

        # TODO: Initialize the required layers (blocks)
        self.layers = nn.ModuleList([ResidualBlock(in_channels, out_channels,
→stride)])
        for i in range(1, num_blocks):
            self.layers.append(ResidualBlock(out_channels, out_channels, 1))

    def forward(self, input):
        # TODO: Execute the layers (blocks)
        #print("Execute " + str(len(self.layers)) + " ResidualBlock layers")
        for layer in self.layers:
            input = layer(input)
        return input

```

Now we are finally ready to implement the full model! To do this, use the `nn.Sequential` API and carefully read the following paragraph from the paper (Fig. 3 is not important):

attachment:resnet_cifar10_description.png

Note that a convolution layer is always convolution + batch norm + activation (ReLU), that each `ResidualBlock` contains 2 layers, and that you might have to **squeeze** the embedding before the dense (fully-connected) layer.

```

[0]: n = 5
      num_classes = 10

      class Squeeze(torch.nn.Module):
          def forward(self, input):
              out = input.squeeze()
              return out

      # TODO: Implement ResNet via nn.Sequential
      resnet = nn.Sequential(
          ResidualStack(3, 16, 1, 1),
          ResidualStack(16, 16, 1, 2 * n),
          ResidualStack(16, 32, 2, 2 * n),
          ResidualStack(32, 64, 2, 2 * n),
          nn.AdaptiveAvgPool2d(1),
          #nn.AdaptiveAvgPool2d(64),
          Squeeze(),
          nn.Linear(64, num_classes)
      )

```

Next we need to initialize the weights of our model.

```
[0]: def initialize_weight(module):
    if isinstance(module, (nn.Linear, nn.Conv2d)):
        nn.init.kaiming_normal_(module.weight, nonlinearity='relu')
    elif isinstance(module, nn.BatchNorm2d):
        nn.init.constant_(module.weight, 1)
        nn.init.constant_(module.bias, 0)

resnet.apply(initialize_weight);
```

5 4. Training

So now we have a shiny new model, but that doesn't really help when we can't train it. So that's what we do next.

First we need to load the data. Note that we split the official training data into train and validation sets, because you must not look at the test set until you are completely done developing your model and report the final results. Some people don't do this properly, but you should not copy other people's bad habits.

```
[0]: class CIFAR10Subset(torchvision.datasets.CIFAR10):
    """
    Get a subset of the CIFAR10 dataset, according to the passed indices.
    """
    def __init__(self, *args, idx=None, **kwargs):
        super().__init__(*args, **kwargs)

        if idx is None:
            return

        self.data = self.data[idx]
        targets_np = np.array(self.targets)
        self.targets = targets_np[idx].tolist()
```

We next define transformations that change the images into PyTorch tensors, standardize the values according to the precomputed mean and standard deviation, and provide data augmentation for the training set.

```
[0]: normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                      std=[0.229, 0.224, 0.225])

transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomCrop(32, 4),
    transforms.ToTensor(),
    normalize,
])

transform_eval = transforms.Compose([
    transforms.ToTensor(),
```

```
        normalize
    ])
```

```
[14]: ntrain = 45_000
train_set = CIFAR10Subset(root='./data', train=True, idx=range(ntrain),
                           download=True, transform=transform_train)
val_set = CIFAR10Subset(root='./data', train=True, idx=range(ntrain, 50_000),
                        download=True, transform=transform_eval)
test_set = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform_eval)
```

0it [00:00, ?it/s]

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
./data/cifar-10-python.tar.gz

170500096it [00:06, 27285630.82it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified
Files already downloaded and verified

```
[0]: dataloaders = {}
dataloaders['train'] = torch.utils.data.DataLoader(train_set, batch_size=128,
                                                    shuffle=True, num_workers=0,
                                                    pin_memory=True)
dataloaders['val'] = torch.utils.data.DataLoader(val_set, batch_size=128,
                                                  shuffle=False, num_workers=0,
                                                  pin_memory=True)
dataloaders['test'] = torch.utils.data.DataLoader(test_set, batch_size=128,
                                                    shuffle=False, num_workers=0,
                                                    pin_memory=True)
```

Next we push the model to our GPU (if there is one).

```
[16]: device = torch.device('cuda') if torch.cuda.is_available() else torch.
      ↪device('cpu')
print(torch.cuda.is_available())
resnet.to(device);
```

True

Next we define a helper method that does one epoch of training or evaluation. We have only defined training here, so you need to implement the necessary changes for evaluation!

```
[0]: def run_epoch(model, optimizer, dataloader, train):
      """
      Run one epoch of training or evaluation.

      Args:
```

```

model: The model used for prediction
optimizer: Optimization algorithm for the model
dataloader: Dataloader providing the data to run our model on
train: Whether this epoch is used for training or evaluation

Returns:
    Loss and accuracy in this epoch.
    """
    # TODO: Change the necessary parts to work correctly during evaluation
    ↪(train=False)

    device = next(model.parameters()).device

    # Set model to training mode (for e.g. batch normalization, dropout)
    model.train()

    epoch_loss = 0.0
    epoch_acc = 0.0

    # Iterate over data
    for xb, yb in dataloader:
        xb, yb = xb.to(device), yb.to(device)

        # zero the parameter gradients
        if train != False:
            optimizer.zero_grad()

        # forward
        with torch.set_grad_enabled(True):
            pred = model(xb)
            loss = F.cross_entropy(pred, yb)
            top1 = torch.argmax(pred, dim=1)
            ncorrect = torch.sum(top1 == yb)

            if train != False:
                loss.backward()
                optimizer.step()

        # statistics
        epoch_loss += loss.item()
        epoch_acc += ncorrect.item()

    epoch_loss /= len(dataloader.dataset)
    epoch_acc /= len(dataloader.dataset)
    return epoch_loss, epoch_acc

```

Next we implement a method for fitting (training) our model. For many models early stopping

can save a lot of training time. Your task is to add early stopping to the loop (based on validation accuracy)! And don't forget to save the best model parameters according to validation accuracy. You will need `copy.deepcopy` and the `state_dict` for this.

```
[0]: def fit(model, optimizer, lr_scheduler, dataloaders, max_epochs, patience):
    """
    Fit the given model on the dataset.

    Args:
        model: The model used for prediction
        optimizer: Optimization algorithm for the model
        lr_scheduler: Learning rate scheduler that improves training
                     in late epochs with learning rate decay
        dataloaders: Dataloaders for training and validation
        max_epochs: Maximum number of epochs for training
        patience: Number of epochs to wait with early stopping the
                 training if validation loss has decreased

    Returns:
        Loss and accuracy in this epoch.
    """

    best_acc = 0.0
    curr_patience = 0
    best_model_weights = None
    for epoch in range(max_epochs):
        train_loss, train_acc = run_epoch(model, optimizer,
        ↪ dataloaders['train'], train=True)
        lr_scheduler.step()
        print(f"Epoch {epoch + 1: >3}/{max_epochs}, train loss: {train_loss:.
        ↪ 2e}, accuracy: {train_acc * 100:.2f}%")

        val_loss, val_acc = run_epoch(model, None, dataloaders['val'],
        ↪ train=False)
        print(f"Epoch {epoch + 1: >3}/{max_epochs}, val loss: {val_loss:.2e},
        ↪ accuracy: {val_acc * 100:.2f}%")

        # TODO: Add early stopping and save the best weights (in
        ↪ best_model_weights)
        curr_patience = 1 if val_acc < best_acc else curr_patience + 1

        updated_model_weights = False
        if best_acc < val_acc:
            best_model_weights = copy.deepcopy(model.state_dict())
            updated_model_weights = True
        best_acc = max(val_acc, best_acc)
```

```

        print(f"Updates: curr_patience={curr_patience}, best_acc={best_acc},  

        ↳updated={updated_model_weights}")

        if curr_patience > patience or epoch + 1 >= max_epochs:
            print(f"Stop early: curr_patience > patience or epoch + 1 >=  

            ↳max_epochs")
            break

    model.load_state_dict(best_model_weights)

```

In most cases you should just use the Adam optimizer for training, because it works well out of the box. However, a well-tuned SGD (with momentum) will in most cases outperform Adam. And since the original paper gives us a well-tuned SGD we will just use that.

```

[19]: optimizer = torch.optim.SGD(resnet.parameters(), lr=0.1, momentum=0.9,  

    ↳weight_decay=1e-4)
    lr_scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, milestones=[100,  

    ↳150], gamma=0.1)

    # Fit model
    fit(resnet, optimizer, lr_scheduler, dataloaders, max_epochs=200, patience=50)

```

```

Epoch 1/200, train loss: 1.80e-02, accuracy: 18.42%
Epoch 1/200, val loss: 1.49e-02, accuracy: 29.16%
Updates: curr_patience=1, best_acc=0.2916, updated=True
Epoch 2/200, train loss: 1.33e-02, accuracy: 36.39%
Epoch 2/200, val loss: 1.22e-02, accuracy: 43.74%
Updates: curr_patience=2, best_acc=0.4374, updated=True
Epoch 3/200, train loss: 1.10e-02, accuracy: 48.72%
Epoch 3/200, val loss: 1.02e-02, accuracy: 56.32%
Updates: curr_patience=3, best_acc=0.5632, updated=True
Epoch 4/200, train loss: 9.33e-03, accuracy: 57.26%
Epoch 4/200, val loss: 8.92e-03, accuracy: 61.78%
Updates: curr_patience=4, best_acc=0.6178, updated=True
Epoch 5/200, train loss: 8.20e-03, accuracy: 62.74%
Epoch 5/200, val loss: 8.05e-03, accuracy: 66.16%
Updates: curr_patience=5, best_acc=0.6616, updated=True
Epoch 6/200, train loss: 7.22e-03, accuracy: 67.73%
Epoch 6/200, val loss: 7.06e-03, accuracy: 70.32%
Updates: curr_patience=6, best_acc=0.7032, updated=True
Epoch 7/200, train loss: 6.42e-03, accuracy: 71.36%
Epoch 7/200, val loss: 6.53e-03, accuracy: 73.50%
Updates: curr_patience=7, best_acc=0.735, updated=True
Epoch 8/200, train loss: 5.82e-03, accuracy: 74.03%
Epoch 8/200, val loss: 6.09e-03, accuracy: 74.92%
Updates: curr_patience=8, best_acc=0.7492, updated=True
Epoch 9/200, train loss: 5.37e-03, accuracy: 76.19%

```

Epoch 9/200, val loss: 5.90e-03, accuracy: 76.60%
 Updates: curr_patience=9, best_acc=0.766, updated=True
 Epoch 10/200, train loss: 4.98e-03, accuracy: 77.90%
 Epoch 10/200, val loss: 5.14e-03, accuracy: 78.54%
 Updates: curr_patience=10, best_acc=0.7854, updated=True
 Epoch 11/200, train loss: 4.67e-03, accuracy: 79.30%
 Epoch 11/200, val loss: 4.94e-03, accuracy: 79.10%
 Updates: curr_patience=11, best_acc=0.791, updated=True
 Epoch 12/200, train loss: 4.41e-03, accuracy: 80.37%
 Epoch 12/200, val loss: 5.12e-03, accuracy: 79.14%
 Updates: curr_patience=12, best_acc=0.7914, updated=True
 Epoch 13/200, train loss: 4.18e-03, accuracy: 81.44%
 Epoch 13/200, val loss: 4.56e-03, accuracy: 80.24%
 Updates: curr_patience=13, best_acc=0.8024, updated=True
 Epoch 14/200, train loss: 3.97e-03, accuracy: 82.44%
 Epoch 14/200, val loss: 4.46e-03, accuracy: 81.68%
 Updates: curr_patience=14, best_acc=0.8168, updated=True
 Epoch 15/200, train loss: 3.85e-03, accuracy: 82.88%
 Epoch 15/200, val loss: 4.50e-03, accuracy: 81.58%
 Updates: curr_patience=1, best_acc=0.8168, updated=False
 Epoch 16/200, train loss: 3.64e-03, accuracy: 83.89%
 Epoch 16/200, val loss: 4.29e-03, accuracy: 82.32%
 Updates: curr_patience=2, best_acc=0.8232, updated=True
 Epoch 17/200, train loss: 3.46e-03, accuracy: 84.62%
 Epoch 17/200, val loss: 4.17e-03, accuracy: 83.06%
 Updates: curr_patience=3, best_acc=0.8306, updated=True
 Epoch 18/200, train loss: 3.39e-03, accuracy: 85.13%
 Epoch 18/200, val loss: 4.26e-03, accuracy: 82.46%
 Updates: curr_patience=1, best_acc=0.8306, updated=False
 Epoch 19/200, train loss: 3.26e-03, accuracy: 85.57%
 Epoch 19/200, val loss: 4.17e-03, accuracy: 83.46%
 Updates: curr_patience=2, best_acc=0.8346, updated=True
 Epoch 20/200, train loss: 3.14e-03, accuracy: 86.09%
 Epoch 20/200, val loss: 4.19e-03, accuracy: 83.56%
 Updates: curr_patience=3, best_acc=0.8356, updated=True
 Epoch 21/200, train loss: 3.08e-03, accuracy: 86.30%
 Epoch 21/200, val loss: 4.05e-03, accuracy: 83.22%
 Updates: curr_patience=1, best_acc=0.8356, updated=False
 Epoch 22/200, train loss: 2.98e-03, accuracy: 86.71%
 Epoch 22/200, val loss: 3.97e-03, accuracy: 83.44%
 Updates: curr_patience=1, best_acc=0.8356, updated=False
 Epoch 23/200, train loss: 2.91e-03, accuracy: 87.10%
 Epoch 23/200, val loss: 4.06e-03, accuracy: 83.42%
 Updates: curr_patience=1, best_acc=0.8356, updated=False
 Epoch 24/200, train loss: 2.83e-03, accuracy: 87.44%
 Epoch 24/200, val loss: 3.85e-03, accuracy: 83.90%
 Updates: curr_patience=2, best_acc=0.839, updated=True
 Epoch 25/200, train loss: 2.79e-03, accuracy: 87.55%

Epoch 25/200, val loss: 3.73e-03, accuracy: 84.46%
Updates: curr_patience=3, best_acc=0.8446, updated=True
Epoch 26/200, train loss: 2.76e-03, accuracy: 87.70%
Epoch 26/200, val loss: 3.97e-03, accuracy: 83.82%
Updates: curr_patience=1, best_acc=0.8446, updated=False
Epoch 27/200, train loss: 2.66e-03, accuracy: 88.26%
Epoch 27/200, val loss: 3.75e-03, accuracy: 84.78%
Updates: curr_patience=2, best_acc=0.8478, updated=True
Epoch 28/200, train loss: 2.59e-03, accuracy: 88.55%
Epoch 28/200, val loss: 3.61e-03, accuracy: 84.90%
Updates: curr_patience=3, best_acc=0.849, updated=True
Epoch 29/200, train loss: 2.52e-03, accuracy: 88.80%
Epoch 29/200, val loss: 3.62e-03, accuracy: 85.06%
Updates: curr_patience=4, best_acc=0.8506, updated=True
Epoch 30/200, train loss: 2.50e-03, accuracy: 88.82%
Epoch 30/200, val loss: 3.58e-03, accuracy: 85.20%
Updates: curr_patience=5, best_acc=0.852, updated=True
Epoch 31/200, train loss: 2.47e-03, accuracy: 89.05%
Epoch 31/200, val loss: 3.59e-03, accuracy: 84.94%
Updates: curr_patience=1, best_acc=0.852, updated=False
Epoch 32/200, train loss: 2.41e-03, accuracy: 89.30%
Epoch 32/200, val loss: 3.64e-03, accuracy: 85.38%
Updates: curr_patience=2, best_acc=0.8538, updated=True
Epoch 33/200, train loss: 2.35e-03, accuracy: 89.53%
Epoch 33/200, val loss: 3.55e-03, accuracy: 86.08%
Updates: curr_patience=3, best_acc=0.8608, updated=True
Epoch 34/200, train loss: 2.34e-03, accuracy: 89.71%
Epoch 34/200, val loss: 3.59e-03, accuracy: 85.30%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 35/200, train loss: 2.31e-03, accuracy: 89.72%
Epoch 35/200, val loss: 3.56e-03, accuracy: 85.54%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 36/200, train loss: 2.25e-03, accuracy: 89.81%
Epoch 36/200, val loss: 3.66e-03, accuracy: 85.32%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 37/200, train loss: 2.22e-03, accuracy: 89.99%
Epoch 37/200, val loss: 3.77e-03, accuracy: 84.86%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 38/200, train loss: 2.26e-03, accuracy: 89.89%
Epoch 38/200, val loss: 3.47e-03, accuracy: 85.88%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 39/200, train loss: 2.13e-03, accuracy: 90.43%
Epoch 39/200, val loss: 3.66e-03, accuracy: 85.62%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 40/200, train loss: 2.13e-03, accuracy: 90.50%
Epoch 40/200, val loss: 3.66e-03, accuracy: 85.08%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 41/200, train loss: 2.10e-03, accuracy: 90.51%

Epoch 41/200, val loss: 3.49e-03, accuracy: 85.82%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 42/200, train loss: 2.07e-03, accuracy: 90.70%
Epoch 42/200, val loss: 3.41e-03, accuracy: 85.86%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 43/200, train loss: 2.04e-03, accuracy: 91.04%
Epoch 43/200, val loss: 3.64e-03, accuracy: 85.92%
Updates: curr_patience=1, best_acc=0.8608, updated=False
Epoch 44/200, train loss: 2.02e-03, accuracy: 90.98%
Epoch 44/200, val loss: 3.52e-03, accuracy: 86.28%
Updates: curr_patience=2, best_acc=0.8628, updated=True
Epoch 45/200, train loss: 1.97e-03, accuracy: 91.14%
Epoch 45/200, val loss: 3.73e-03, accuracy: 85.34%
Updates: curr_patience=1, best_acc=0.8628, updated=False
Epoch 46/200, train loss: 1.95e-03, accuracy: 91.30%
Epoch 46/200, val loss: 3.82e-03, accuracy: 85.20%
Updates: curr_patience=1, best_acc=0.8628, updated=False
Epoch 47/200, train loss: 1.98e-03, accuracy: 91.17%
Epoch 47/200, val loss: 3.38e-03, accuracy: 86.98%
Updates: curr_patience=2, best_acc=0.8698, updated=True
Epoch 48/200, train loss: 1.93e-03, accuracy: 91.34%
Epoch 48/200, val loss: 3.77e-03, accuracy: 85.46%
Updates: curr_patience=1, best_acc=0.8698, updated=False
Epoch 49/200, train loss: 1.91e-03, accuracy: 91.52%
Epoch 49/200, val loss: 3.58e-03, accuracy: 85.54%
Updates: curr_patience=1, best_acc=0.8698, updated=False
Epoch 50/200, train loss: 1.87e-03, accuracy: 91.49%
Epoch 50/200, val loss: 3.35e-03, accuracy: 86.56%
Updates: curr_patience=1, best_acc=0.8698, updated=False
Epoch 51/200, train loss: 1.86e-03, accuracy: 91.70%
Epoch 51/200, val loss: 3.23e-03, accuracy: 87.06%
Updates: curr_patience=2, best_acc=0.8706, updated=True
Epoch 52/200, train loss: 1.89e-03, accuracy: 91.59%
Epoch 52/200, val loss: 3.36e-03, accuracy: 87.02%
Updates: curr_patience=1, best_acc=0.8706, updated=False
Epoch 53/200, train loss: 1.82e-03, accuracy: 91.88%
Epoch 53/200, val loss: 3.52e-03, accuracy: 86.30%
Updates: curr_patience=1, best_acc=0.8706, updated=False
Epoch 54/200, train loss: 1.81e-03, accuracy: 92.00%
Epoch 54/200, val loss: 3.50e-03, accuracy: 86.86%
Updates: curr_patience=1, best_acc=0.8706, updated=False
Epoch 55/200, train loss: 1.81e-03, accuracy: 91.93%
Epoch 55/200, val loss: 3.45e-03, accuracy: 86.64%
Updates: curr_patience=1, best_acc=0.8706, updated=False
Epoch 56/200, train loss: 1.80e-03, accuracy: 91.83%
Epoch 56/200, val loss: 3.57e-03, accuracy: 86.52%
Updates: curr_patience=1, best_acc=0.8706, updated=False
Epoch 57/200, train loss: 1.72e-03, accuracy: 92.24%

Epoch 57/200, val loss: 3.13e-03, accuracy: 87.72%
Updates: curr_patience=2, best_acc=0.8772, updated=True
Epoch 58/200, train loss: 1.74e-03, accuracy: 92.21%
Epoch 58/200, val loss: 3.41e-03, accuracy: 86.68%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 59/200, train loss: 1.74e-03, accuracy: 92.22%
Epoch 59/200, val loss: 3.19e-03, accuracy: 87.32%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 60/200, train loss: 1.69e-03, accuracy: 92.51%
Epoch 60/200, val loss: 3.34e-03, accuracy: 86.86%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 61/200, train loss: 1.71e-03, accuracy: 92.39%
Epoch 61/200, val loss: 3.46e-03, accuracy: 87.16%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 62/200, train loss: 1.71e-03, accuracy: 92.21%
Epoch 62/200, val loss: 3.41e-03, accuracy: 86.38%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 63/200, train loss: 1.71e-03, accuracy: 92.17%
Epoch 63/200, val loss: 3.24e-03, accuracy: 86.92%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 64/200, train loss: 1.67e-03, accuracy: 92.52%
Epoch 64/200, val loss: 3.53e-03, accuracy: 86.58%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 65/200, train loss: 1.65e-03, accuracy: 92.71%
Epoch 65/200, val loss: 3.40e-03, accuracy: 87.32%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 66/200, train loss: 1.67e-03, accuracy: 92.53%
Epoch 66/200, val loss: 3.47e-03, accuracy: 87.02%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 67/200, train loss: 1.59e-03, accuracy: 92.79%
Epoch 67/200, val loss: 3.63e-03, accuracy: 86.06%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 68/200, train loss: 1.69e-03, accuracy: 92.32%
Epoch 68/200, val loss: 3.37e-03, accuracy: 87.30%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 69/200, train loss: 1.58e-03, accuracy: 92.79%
Epoch 69/200, val loss: 3.41e-03, accuracy: 86.68%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 70/200, train loss: 1.61e-03, accuracy: 92.82%
Epoch 70/200, val loss: 3.25e-03, accuracy: 87.28%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 71/200, train loss: 1.57e-03, accuracy: 92.92%
Epoch 71/200, val loss: 3.47e-03, accuracy: 86.84%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 72/200, train loss: 1.61e-03, accuracy: 92.76%
Epoch 72/200, val loss: 3.41e-03, accuracy: 86.86%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 73/200, train loss: 1.57e-03, accuracy: 92.85%

Epoch 73/200, val loss: 3.42e-03, accuracy: 87.22%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 74/200, train loss: 1.56e-03, accuracy: 92.99%
Epoch 74/200, val loss: 3.34e-03, accuracy: 87.14%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 75/200, train loss: 1.52e-03, accuracy: 93.29%
Epoch 75/200, val loss: 3.32e-03, accuracy: 87.04%
Updates: curr_patience=1, best_acc=0.8772, updated=False
Epoch 76/200, train loss: 1.48e-03, accuracy: 93.42%
Epoch 76/200, val loss: 3.25e-03, accuracy: 88.00%
Updates: curr_patience=2, best_acc=0.88, updated=True
Epoch 77/200, train loss: 1.56e-03, accuracy: 93.10%
Epoch 77/200, val loss: 3.36e-03, accuracy: 86.92%
Updates: curr_patience=1, best_acc=0.88, updated=False
Epoch 78/200, train loss: 1.52e-03, accuracy: 93.17%
Epoch 78/200, val loss: 3.47e-03, accuracy: 86.96%
Updates: curr_patience=1, best_acc=0.88, updated=False
Epoch 79/200, train loss: 1.51e-03, accuracy: 93.28%
Epoch 79/200, val loss: 3.39e-03, accuracy: 86.76%
Updates: curr_patience=1, best_acc=0.88, updated=False
Epoch 80/200, train loss: 1.51e-03, accuracy: 93.23%
Epoch 80/200, val loss: 2.96e-03, accuracy: 88.20%
Updates: curr_patience=2, best_acc=0.882, updated=True
Epoch 81/200, train loss: 1.51e-03, accuracy: 93.16%
Epoch 81/200, val loss: 3.22e-03, accuracy: 87.18%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 82/200, train loss: 1.47e-03, accuracy: 93.48%
Epoch 82/200, val loss: 3.16e-03, accuracy: 88.14%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 83/200, train loss: 1.45e-03, accuracy: 93.40%
Epoch 83/200, val loss: 3.43e-03, accuracy: 86.44%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 84/200, train loss: 1.51e-03, accuracy: 93.24%
Epoch 84/200, val loss: 3.21e-03, accuracy: 87.50%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 85/200, train loss: 1.47e-03, accuracy: 93.30%
Epoch 85/200, val loss: 3.37e-03, accuracy: 87.58%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 86/200, train loss: 1.50e-03, accuracy: 93.12%
Epoch 86/200, val loss: 3.38e-03, accuracy: 87.70%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 87/200, train loss: 1.49e-03, accuracy: 93.28%
Epoch 87/200, val loss: 3.39e-03, accuracy: 87.30%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 88/200, train loss: 1.43e-03, accuracy: 93.54%
Epoch 88/200, val loss: 3.27e-03, accuracy: 87.44%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 89/200, train loss: 1.50e-03, accuracy: 93.20%

Epoch 89/200, val loss: 3.42e-03, accuracy: 87.02%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 90/200, train loss: 1.48e-03, accuracy: 93.43%
Epoch 90/200, val loss: 3.18e-03, accuracy: 88.12%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 91/200, train loss: 1.47e-03, accuracy: 93.41%
Epoch 91/200, val loss: 3.07e-03, accuracy: 88.04%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 92/200, train loss: 1.43e-03, accuracy: 93.60%
Epoch 92/200, val loss: 3.27e-03, accuracy: 87.80%
Updates: curr_patience=1, best_acc=0.882, updated=False
Epoch 93/200, train loss: 1.45e-03, accuracy: 93.43%
Epoch 93/200, val loss: 3.30e-03, accuracy: 88.28%
Updates: curr_patience=2, best_acc=0.8828, updated=True
Epoch 94/200, train loss: 1.43e-03, accuracy: 93.68%
Epoch 94/200, val loss: 3.32e-03, accuracy: 87.36%
Updates: curr_patience=1, best_acc=0.8828, updated=False
Epoch 95/200, train loss: 1.42e-03, accuracy: 93.63%
Epoch 95/200, val loss: 3.28e-03, accuracy: 88.18%
Updates: curr_patience=1, best_acc=0.8828, updated=False
Epoch 96/200, train loss: 1.43e-03, accuracy: 93.50%
Epoch 96/200, val loss: 3.59e-03, accuracy: 87.06%
Updates: curr_patience=1, best_acc=0.8828, updated=False
Epoch 97/200, train loss: 1.43e-03, accuracy: 93.69%
Epoch 97/200, val loss: 3.36e-03, accuracy: 87.60%
Updates: curr_patience=1, best_acc=0.8828, updated=False
Epoch 98/200, train loss: 1.45e-03, accuracy: 93.45%
Epoch 98/200, val loss: 3.55e-03, accuracy: 86.86%
Updates: curr_patience=1, best_acc=0.8828, updated=False
Epoch 99/200, train loss: 1.39e-03, accuracy: 93.53%
Epoch 99/200, val loss: 3.01e-03, accuracy: 88.44%
Updates: curr_patience=2, best_acc=0.8844, updated=True
Epoch 100/200, train loss: 1.38e-03, accuracy: 93.80%
Epoch 100/200, val loss: 3.61e-03, accuracy: 86.70%
Updates: curr_patience=1, best_acc=0.8844, updated=False
Epoch 101/200, train loss: 7.20e-04, accuracy: 96.86%
Epoch 101/200, val loss: 2.71e-03, accuracy: 89.98%
Updates: curr_patience=2, best_acc=0.8998, updated=True
Epoch 102/200, train loss: 4.79e-04, accuracy: 98.01%
Epoch 102/200, val loss: 2.67e-03, accuracy: 90.60%
Updates: curr_patience=3, best_acc=0.906, updated=True
Epoch 103/200, train loss: 4.08e-04, accuracy: 98.33%
Epoch 103/200, val loss: 2.69e-03, accuracy: 90.68%
Updates: curr_patience=4, best_acc=0.9068, updated=True
Epoch 104/200, train loss: 3.61e-04, accuracy: 98.54%
Epoch 104/200, val loss: 2.77e-03, accuracy: 90.70%
Updates: curr_patience=5, best_acc=0.907, updated=True
Epoch 105/200, train loss: 3.31e-04, accuracy: 98.62%

Epoch 105/200, val loss: 2.77e-03, accuracy: 90.70%
Updates: curr_patience=6, best_acc=0.907, updated=False
Epoch 106/200, train loss: 2.93e-04, accuracy: 98.82%
Epoch 106/200, val loss: 2.89e-03, accuracy: 90.58%
Updates: curr_patience=1, best_acc=0.907, updated=False
Epoch 107/200, train loss: 2.73e-04, accuracy: 98.89%
Epoch 107/200, val loss: 2.88e-03, accuracy: 90.98%
Updates: curr_patience=2, best_acc=0.9098, updated=True
Epoch 108/200, train loss: 2.46e-04, accuracy: 98.98%
Epoch 108/200, val loss: 2.92e-03, accuracy: 90.92%
Updates: curr_patience=1, best_acc=0.9098, updated=False
Epoch 109/200, train loss: 2.32e-04, accuracy: 99.04%
Epoch 109/200, val loss: 2.95e-03, accuracy: 90.86%
Updates: curr_patience=1, best_acc=0.9098, updated=False
Epoch 110/200, train loss: 1.99e-04, accuracy: 99.18%
Epoch 110/200, val loss: 3.06e-03, accuracy: 90.78%
Updates: curr_patience=1, best_acc=0.9098, updated=False
Epoch 111/200, train loss: 1.95e-04, accuracy: 99.27%
Epoch 111/200, val loss: 3.08e-03, accuracy: 90.64%
Updates: curr_patience=1, best_acc=0.9098, updated=False
Epoch 112/200, train loss: 1.89e-04, accuracy: 99.27%
Epoch 112/200, val loss: 3.04e-03, accuracy: 91.00%
Updates: curr_patience=2, best_acc=0.91, updated=True
Epoch 113/200, train loss: 1.70e-04, accuracy: 99.38%
Epoch 113/200, val loss: 3.11e-03, accuracy: 90.62%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 114/200, train loss: 1.74e-04, accuracy: 99.31%
Epoch 114/200, val loss: 3.11e-03, accuracy: 90.70%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 115/200, train loss: 1.69e-04, accuracy: 99.33%
Epoch 115/200, val loss: 3.19e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 116/200, train loss: 1.52e-04, accuracy: 99.42%
Epoch 116/200, val loss: 3.29e-03, accuracy: 90.48%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 117/200, train loss: 1.43e-04, accuracy: 99.44%
Epoch 117/200, val loss: 3.37e-03, accuracy: 90.54%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 118/200, train loss: 1.37e-04, accuracy: 99.48%
Epoch 118/200, val loss: 3.31e-03, accuracy: 90.36%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 119/200, train loss: 1.31e-04, accuracy: 99.48%
Epoch 119/200, val loss: 3.29e-03, accuracy: 90.68%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 120/200, train loss: 1.27e-04, accuracy: 99.50%
Epoch 120/200, val loss: 3.30e-03, accuracy: 90.62%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 121/200, train loss: 1.31e-04, accuracy: 99.47%

Epoch 121/200, val loss: 3.48e-03, accuracy: 90.50%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 122/200, train loss: 1.12e-04, accuracy: 99.60%
 Epoch 122/200, val loss: 3.40e-03, accuracy: 90.68%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 123/200, train loss: 1.15e-04, accuracy: 99.52%
 Epoch 123/200, val loss: 3.55e-03, accuracy: 90.34%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 124/200, train loss: 1.09e-04, accuracy: 99.59%
 Epoch 124/200, val loss: 3.42e-03, accuracy: 90.62%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 125/200, train loss: 1.06e-04, accuracy: 99.59%
 Epoch 125/200, val loss: 3.45e-03, accuracy: 90.58%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 126/200, train loss: 9.50e-05, accuracy: 99.66%
 Epoch 126/200, val loss: 3.41e-03, accuracy: 90.88%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 127/200, train loss: 1.05e-04, accuracy: 99.64%
 Epoch 127/200, val loss: 3.53e-03, accuracy: 90.42%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 128/200, train loss: 9.98e-05, accuracy: 99.62%
 Epoch 128/200, val loss: 3.42e-03, accuracy: 90.54%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 129/200, train loss: 1.00e-04, accuracy: 99.58%
 Epoch 129/200, val loss: 3.55e-03, accuracy: 90.66%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 130/200, train loss: 8.87e-05, accuracy: 99.68%
 Epoch 130/200, val loss: 3.49e-03, accuracy: 90.94%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 131/200, train loss: 8.15e-05, accuracy: 99.73%
 Epoch 131/200, val loss: 3.52e-03, accuracy: 90.76%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 132/200, train loss: 8.20e-05, accuracy: 99.70%
 Epoch 132/200, val loss: 3.52e-03, accuracy: 90.74%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 133/200, train loss: 7.63e-05, accuracy: 99.74%
 Epoch 133/200, val loss: 3.59e-03, accuracy: 90.70%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 134/200, train loss: 8.69e-05, accuracy: 99.69%
 Epoch 134/200, val loss: 3.54e-03, accuracy: 90.84%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 135/200, train loss: 7.38e-05, accuracy: 99.76%
 Epoch 135/200, val loss: 3.54e-03, accuracy: 90.74%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 136/200, train loss: 7.20e-05, accuracy: 99.72%
 Epoch 136/200, val loss: 3.64e-03, accuracy: 90.50%
 Updates: curr_patience=1, best_acc=0.91, updated=False
 Epoch 137/200, train loss: 7.26e-05, accuracy: 99.76%

Epoch 137/200, val loss: 3.79e-03, accuracy: 90.32%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 138/200, train loss: 7.19e-05, accuracy: 99.76%
Epoch 138/200, val loss: 3.56e-03, accuracy: 90.62%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 139/200, train loss: 7.23e-05, accuracy: 99.74%
Epoch 139/200, val loss: 3.60e-03, accuracy: 90.44%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 140/200, train loss: 7.39e-05, accuracy: 99.74%
Epoch 140/200, val loss: 3.68e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.91, updated=False
Epoch 141/200, train loss: 7.02e-05, accuracy: 99.77%
Epoch 141/200, val loss: 3.74e-03, accuracy: 91.10%
Updates: curr_patience=2, best_acc=0.911, updated=True
Epoch 142/200, train loss: 6.90e-05, accuracy: 99.75%
Epoch 142/200, val loss: 3.73e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 143/200, train loss: 6.73e-05, accuracy: 99.77%
Epoch 143/200, val loss: 3.74e-03, accuracy: 90.86%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 144/200, train loss: 6.83e-05, accuracy: 99.78%
Epoch 144/200, val loss: 3.70e-03, accuracy: 90.84%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 145/200, train loss: 6.36e-05, accuracy: 99.78%
Epoch 145/200, val loss: 3.64e-03, accuracy: 91.04%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 146/200, train loss: 6.35e-05, accuracy: 99.78%
Epoch 146/200, val loss: 3.82e-03, accuracy: 90.70%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 147/200, train loss: 6.96e-05, accuracy: 99.74%
Epoch 147/200, val loss: 3.80e-03, accuracy: 90.50%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 148/200, train loss: 6.83e-05, accuracy: 99.74%
Epoch 148/200, val loss: 3.73e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 149/200, train loss: 5.86e-05, accuracy: 99.80%
Epoch 149/200, val loss: 3.81e-03, accuracy: 90.68%
Updates: curr_patience=1, best_acc=0.911, updated=False
Epoch 150/200, train loss: 5.65e-05, accuracy: 99.79%
Epoch 150/200, val loss: 3.82e-03, accuracy: 91.18%
Updates: curr_patience=2, best_acc=0.9118, updated=True
Epoch 151/200, train loss: 5.20e-05, accuracy: 99.80%
Epoch 151/200, val loss: 3.79e-03, accuracy: 91.14%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 152/200, train loss: 4.47e-05, accuracy: 99.86%
Epoch 152/200, val loss: 3.79e-03, accuracy: 91.16%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 153/200, train loss: 4.52e-05, accuracy: 99.86%

Epoch 153/200, val loss: 3.78e-03, accuracy: 91.16%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 154/200, train loss: 4.77e-05, accuracy: 99.84%
Epoch 154/200, val loss: 3.80e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 155/200, train loss: 4.24e-05, accuracy: 99.87%
Epoch 155/200, val loss: 3.79e-03, accuracy: 90.98%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 156/200, train loss: 4.05e-05, accuracy: 99.87%
Epoch 156/200, val loss: 3.79e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 157/200, train loss: 3.85e-05, accuracy: 99.91%
Epoch 157/200, val loss: 3.79e-03, accuracy: 90.92%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 158/200, train loss: 3.74e-05, accuracy: 99.88%
Epoch 158/200, val loss: 3.78e-03, accuracy: 90.88%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 159/200, train loss: 3.89e-05, accuracy: 99.89%
Epoch 159/200, val loss: 3.78e-03, accuracy: 90.92%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 160/200, train loss: 3.73e-05, accuracy: 99.91%
Epoch 160/200, val loss: 3.79e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 161/200, train loss: 3.73e-05, accuracy: 99.90%
Epoch 161/200, val loss: 3.81e-03, accuracy: 90.86%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 162/200, train loss: 3.79e-05, accuracy: 99.88%
Epoch 162/200, val loss: 3.79e-03, accuracy: 90.92%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 163/200, train loss: 3.51e-05, accuracy: 99.92%
Epoch 163/200, val loss: 3.79e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 164/200, train loss: 3.81e-05, accuracy: 99.90%
Epoch 164/200, val loss: 3.78e-03, accuracy: 91.00%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 165/200, train loss: 3.74e-05, accuracy: 99.89%
Epoch 165/200, val loss: 3.77e-03, accuracy: 91.10%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 166/200, train loss: 3.34e-05, accuracy: 99.90%
Epoch 166/200, val loss: 3.78e-03, accuracy: 91.04%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 167/200, train loss: 3.06e-05, accuracy: 99.93%
Epoch 167/200, val loss: 3.80e-03, accuracy: 91.10%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 168/200, train loss: 3.33e-05, accuracy: 99.91%
Epoch 168/200, val loss: 3.80e-03, accuracy: 91.04%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 169/200, train loss: 3.57e-05, accuracy: 99.89%

Epoch 169/200, val loss: 3.78e-03, accuracy: 91.14%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 170/200, train loss: 3.24e-05, accuracy: 99.91%
Epoch 170/200, val loss: 3.80e-03, accuracy: 91.12%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 171/200, train loss: 3.44e-05, accuracy: 99.91%
Epoch 171/200, val loss: 3.80e-03, accuracy: 91.02%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 172/200, train loss: 3.54e-05, accuracy: 99.90%
Epoch 172/200, val loss: 3.80e-03, accuracy: 91.08%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 173/200, train loss: 3.16e-05, accuracy: 99.92%
Epoch 173/200, val loss: 3.81e-03, accuracy: 91.00%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 174/200, train loss: 3.56e-05, accuracy: 99.90%
Epoch 174/200, val loss: 3.80e-03, accuracy: 90.88%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 175/200, train loss: 3.16e-05, accuracy: 99.92%
Epoch 175/200, val loss: 3.81e-03, accuracy: 90.88%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 176/200, train loss: 3.28e-05, accuracy: 99.90%
Epoch 176/200, val loss: 3.80e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 177/200, train loss: 3.20e-05, accuracy: 99.91%
Epoch 177/200, val loss: 3.81e-03, accuracy: 90.88%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 178/200, train loss: 3.41e-05, accuracy: 99.89%
Epoch 178/200, val loss: 3.81e-03, accuracy: 90.96%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 179/200, train loss: 3.32e-05, accuracy: 99.90%
Epoch 179/200, val loss: 3.81e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 180/200, train loss: 3.01e-05, accuracy: 99.93%
Epoch 180/200, val loss: 3.81e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 181/200, train loss: 3.18e-05, accuracy: 99.92%
Epoch 181/200, val loss: 3.82e-03, accuracy: 90.78%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 182/200, train loss: 2.84e-05, accuracy: 99.93%
Epoch 182/200, val loss: 3.82e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 183/200, train loss: 3.19e-05, accuracy: 99.92%
Epoch 183/200, val loss: 3.81e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 184/200, train loss: 3.21e-05, accuracy: 99.92%
Epoch 184/200, val loss: 3.81e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 185/200, train loss: 3.09e-05, accuracy: 99.90%

Epoch 185/200, val loss: 3.82e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 186/200, train loss: 3.11e-05, accuracy: 99.92%
Epoch 186/200, val loss: 3.82e-03, accuracy: 90.84%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 187/200, train loss: 3.06e-05, accuracy: 99.92%
Epoch 187/200, val loss: 3.82e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 188/200, train loss: 2.91e-05, accuracy: 99.94%
Epoch 188/200, val loss: 3.82e-03, accuracy: 90.86%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 189/200, train loss: 3.08e-05, accuracy: 99.94%
Epoch 189/200, val loss: 3.83e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 190/200, train loss: 3.01e-05, accuracy: 99.91%
Epoch 190/200, val loss: 3.81e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 191/200, train loss: 2.96e-05, accuracy: 99.93%
Epoch 191/200, val loss: 3.81e-03, accuracy: 90.94%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 192/200, train loss: 3.31e-05, accuracy: 99.91%
Epoch 192/200, val loss: 3.82e-03, accuracy: 90.86%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 193/200, train loss: 2.78e-05, accuracy: 99.93%
Epoch 193/200, val loss: 3.82e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 194/200, train loss: 2.98e-05, accuracy: 99.92%
Epoch 194/200, val loss: 3.83e-03, accuracy: 90.72%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 195/200, train loss: 2.58e-05, accuracy: 99.94%
Epoch 195/200, val loss: 3.83e-03, accuracy: 90.76%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 196/200, train loss: 2.89e-05, accuracy: 99.93%
Epoch 196/200, val loss: 3.82e-03, accuracy: 90.84%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 197/200, train loss: 2.61e-05, accuracy: 99.94%
Epoch 197/200, val loss: 3.83e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 198/200, train loss: 2.73e-05, accuracy: 99.94%
Epoch 198/200, val loss: 3.84e-03, accuracy: 90.80%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 199/200, train loss: 2.86e-05, accuracy: 99.93%
Epoch 199/200, val loss: 3.83e-03, accuracy: 90.82%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Epoch 200/200, train loss: 3.48e-05, accuracy: 99.90%
Epoch 200/200, val loss: 3.82e-03, accuracy: 90.88%
Updates: curr_patience=1, best_acc=0.9118, updated=False
Stop early: curr_patience > patience or epoch + 1 >= max_epochs

Once the model is trained we run it on the test set to obtain our final accuracy. Note that we can only look at the test set once, everything else would lead to overfitting. So you *must* ignore the test set while developing your model!

```
[20]: test_loss, test_acc = run_epoch(resnet, None, dataloaders['test'], train=False)
      print(f"Test loss: {test_loss:.1e}, accuracy: {test_acc * 100:.2f}%")
```

Test loss: 3.6e-03, accuracy: 90.36%

That's almost what was reported in the paper (92.49%) and we didn't even train on the full training set.

6 Optional task: Squeeze out all the juice!

Can you do even better? Have a look at [A Recipe for Training Neural Networks](#) and at the [EfficientNet architecture](#) we discussed in the lecture. Play around with the possibilities PyTorch offers you and see how close you can get to the [state of the art on CIFAR-10](#).

Hint: You can use [Google Colab](#) to access some free GPUs for your experiments.

Appendix

We confirm that the submitted solution is original work and was written by us without further assistance.
Appropriate credit has been given where reference has been made to the work of others.

Munich, January 2, 2020, Signature Marcel Bruckner (03674122)

Munich, January 2, 2020, Signature Julian Hohenadel (03673879)

Munich, January 2, 2020, Signature Kevin Bein (03707775)