exercise 10 notebook 90 36

January 5, 2020

```
import numpy as np
   import torch
   import torch.nn as nn
   import torch.nn.functional as F
   import torchvision
   import torchvision.transforms as transforms
[2]: !nvidia-smi
   Fri Jan 3 19:33:22 2020
   | NVIDIA-SMI 440.44
                    Driver Version: 418.67 CUDA Version: 10.1
   |-----
   | GPU Name
                Persistence-M | Bus-Id
                                    Disp.A | Volatile Uncorr. ECC |
   | Fan Temp Perf Pwr:Usage/Cap| Memory-Usage | GPU-Util Compute M. |
   |------
     O Tesla P100-PCIE... Off | 00000000:00:04.0 Off |
                                                       0 |
                              OMiB / 16280MiB |
           PO 28W / 250W |
                                                    Default |
   | Processes:
                                                  GPU Memory
            PID
                                                  Usage
                Type Process name
   |-----|
    No running processes found
```

1 PyTorch

[0]: import copy

In this notebook you will gain some hands-on experience with PyTorch, one of the major frameworks for deep learning. To install PyTorch run conda install pytorch torchvision cudatoolkit=10.1 -c pytorch, with cudatoolkit set to whichever CUDA version you have installed. You can check this by running nvcc --version. If you do not have an Nvidia GPU you can run conda install pytorch torchvision cpuonly -c pytorch instead. However, in this case we recommend using Google Colab.

You will start by re-implementing some common features of deep neural networks (dropout and

batch normalization) and then implement a very popular modern architecture for image classification (ResNet) and improve its training loop.

2 1. Dropout

Dropout is a form of regularization for neural networks. It works by randomly setting activations (values) to 0, each one with equal probability p. The values are then scaled by a factor $\frac{1}{1-p}$ to conserve their mean.

Dropout effectively trains a pseudo-ensemble of models with stochastic gradient descent. During evaluation we want to use the full ensemble and therefore have to turn off dropout. Use self.training to check if the model is in training or evaluation mode.

```
[0]: class Dropout(nn.Module):
         11 II II
         Dropout, as discussed in the lecture and described here:
         https://pytorch.org/docs/stable/nn.html#torch.nn.Dropout
             p: float, dropout probability
         def __init__(self, p):
             super().__init__()
             self.p = p
         def forward(self, input):
             The module's forward pass.
             This has to be implemented for every PyTorch module.
             PyTorch then automatically generates the backward pass
             by dynamically generating the computational graph during
             execution.
             Args:
                 input: PyTorch tensor, arbitrary shape
             Returns:
                 PyTorch tensor, same shape as input
             if self.training:
                 mask = np.random.random(input.shape)
                 return torch.from_numpy(np.where(mask <= self.p, 0, (1 / (1 - self.
      →p))))
             # TODO: Set values randomly to 0.
```

```
[4]: # Test dropout
    test = torch.ones(10_000)
    dropout = Dropout(0.5)
    test_dropped = dropout(test)

print(test_dropped.sum().item())
print((test_dropped > 0).sum().item())

print(np.isclose(test_dropped.sum().item(), 10_000, atol=400))
print(np.isclose((test_dropped > 0).sum().item(), 5_000, atol=200))

# These assertions can in principle fail due to bad luck, but
# if implemented correctly they should almost always succeed.
assert np.isclose(test_dropped.sum().item(), 10_000, atol=400)
assert np.isclose((test_dropped > 0).sum().item(), 5_000, atol=200)
```

10104.0 5052 True True

3 2. Batch normalization

Batch normalization is a trick used to smoothen the loss landscape and improve training. It is defined as the function

$$y = \frac{x - \mu_x}{\sigma_x + \epsilon} \cdot \gamma + \beta$$

, where γ and β and learnable parameters and ϵ is a some small number to avoid dividing by zero. The Statistics μ_x and σ_x are taken separately for each feature. In a CNN this means averaging over the batch and all pixels.

```
[0]: class BatchNorm(nn.Module):

"""

Batch normalization, as discussed in the lecture and similar to https://pytorch.org/docs/stable/nn.html#torch.nn.BatchNorm1d

Only uses batch statistics (no running mean for evaluation).

Batch statistics are calculated for a single dimension.

Gamma is initialized as 1, beta as 0.

Args:

num_features: Number of features to calculate batch statistics for.

"""

def __init__(self, num_features):
    super().__init__()

# TODO: Initialize the required parameters
```

```
self.gamma = nn.Parameter(torch.ones(num_features)).unsqueeze(0).

unsqueeze(-1)

       \#self.gamma = self.gamma.
       self.beta = nn.Parameter(torch.zeros(num_features)).unsqueeze(0).
\rightarrowunsqueeze(-1)
       #self.beta = self.beta.
  def forward(self, input):
       Batch normalization over the dimension C of (N, C, L).
       Args:
           input: PyTorch tensor, shape [N, C, L]
           PyTorch tensor, same shape as input
       eps = 1e-5
       mean = input.mean(dim=[0, 2], keepdim=True)
       input_mean_norm = input - mean
       var = torch.sqrt(input.var(dim=[0, 2], keepdim=True))
       return (input_mean_norm / (var + eps)) * self.gamma + self.beta
       # TODO: Implement the required transformation
```

```
[6]: # Tests the batch normalization implementation
    torch.random.manual_seed(42)
    test = torch.randn(8, 2, 4)

b1 = BatchNorm(2)
    test_b1 = b1(test)

b2 = nn.BatchNorm1d(2, affine=False, track_running_stats=False)
    test_b2 = b2(test)

print(test_b1)
    print("----")
    print(test_b2)

print(torch.allclose(test_b1, test_b2, rtol=0.02))
    assert torch.allclose(test_b1, test_b2, rtol=0.02)
```

```
tensor([[[ 1.6380, 1.2470, 0.7253, -1.9484], [ 0.6971, -1.2133, -0.0234, -1.5829]],
```

```
[[-0.7447, 1.3905, -0.4249, -1.3242],
        [-0.7073, -0.5391, -0.7482, 0.7810]],
        [[1.3849, -0.2178, -0.5182, 0.3152],
        [-0.7375, 1.0964, 0.8193, 1.6979]],
        [[1.0618, 1.0772, 0.4671, 1.1113],
        [-0.2117, 0.0613, -0.2317, 0.8782]],
        [[-1.3073, -0.8507, -0.2745, 1.4516],
        [0.3380, -0.4044, 0.3249, -0.7540]],
        [[-1.4611, 0.8097, -0.8583, -0.6105],
        [-1.2529, 2.1395, -1.2134, -0.4677]],
        [[-0.8886, -0.6611, -0.0064, 0.3918],
        [-0.4678, 1.2093, -0.7933, -0.7154]],
        [[-1.3238, -0.0438, -0.1323, 0.5251],
        [-0.0781, 1.8616, -1.1634, 1.4012]]], grad_fn=<AddBackward0>)
tensor([[[ 1.6642, 1.2669, 0.7369, -1.9796],
        [0.7082, -1.2327, -0.0238, -1.6083]],
        [[-0.7567, 1.4128, -0.4317, -1.3454],
        [-0.7186, -0.5477, -0.7602, 0.7935]],
        [[1.4070, -0.2212, -0.5265, 0.3202],
        [-0.7493, 1.1140, 0.8324, 1.7251]],
        [[ 1.0788, 1.0945, 0.4746, 1.1291],
        [-0.2151, 0.0622, -0.2354, 0.8923]],
        [[-1.3282, -0.8643, -0.2789, 1.4748],
        [0.3434, -0.4109, 0.3301, -0.7660]],
        [[-1.4845, 0.8227, -0.8720, -0.6202],
        [-1.2729, 2.1737, -1.2328, -0.4752]],
        [[-0.9028, -0.6717, -0.0065, 0.3981],
        [-0.4753, 1.2287, -0.8060, -0.7269]],
        [[-1.3450, -0.0445, -0.1344, 0.5335],
        [-0.0794, 1.8914, -1.1820, 1.4237]]])
True
```

4 3. ResNet

ResNet is the model that first introduced residual connections (a form of skip connections). It is a rather simple, but successful and very popular architecture. In this part of the exercise we will re-implement it step by step.

Note that there is also an improved version of ResNet with optimized residual blocks. Here we will implement the original version for CIFAR-10.

This is just a convenience function to make e.g. nn.Sequential more flexible. It is e.g. useful in combination with x.squeeze().

```
[0]: class Lambda(nn.Module):
    def __init__(self, func):
        super().__init__()
        self.func = func

    def forward(self, x):
        return self.func(x)
```

We begin by implementing the residual blocks. The block is illustrated by this sketch:

attachment:residual_connection.png

Note that we use 'SAME' padding, no bias, and batch normalization after each convolution.

```
[0]: class ResidualBlock(nn.Module):
         11 11 11
         The residual block used by ResNet.
         Args:
             in_channels: The number of channels (feature maps) of the incoming_
      \hookrightarrow embedding
             out_channels: The number of channels after the first convolution
             stride: Stride size of the first convolution, used for downsampling
         def __init__(self, in_channels, out_channels, stride=1):
             super().__init__()
             if stride > 1 or in_channels != out_channels:
                 # Add strides in the skip connection and zeros for the new channels.
                 self.skip = Lambda(lambda x: F.pad(x[:, :, ::stride, ::stride],
                                                      (0, 0, 0, 0, 0 out_channels -u
      →in_channels),
                                                      mode="constant", value=0))
             else:
                 self.skip = nn.Sequential()
             # TODO: Initialize the required layers
```

```
#self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,_
\rightarrow padding="SAME", bias=False, stride=stride)
       self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,__
→padding=1, bias=False, stride=stride, groups=1)
       self.bn1 = nn.BatchNorm2d(out_channels) #BatchNorm(out_channels)
       #self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,_
\rightarrow padding="SAME", bias=False, stride=stride)
       self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,_
→padding=1, bias=False, stride=1, groups=1)
       self.bn2 = nn.BatchNorm2d(out_channels)
       self.relu = nn.ReLU()
  def forward(self, input):
       # TODO: Execute the required layers and functions
       #print("input.shape:", input.shape)
       identity = self.skip(input)
       #print("identity.shape:", identity.shape)
       out = self.conv1(input)
       out = self.bn1(out)
       out = self.relu(out)
       #print("out1.shape:", out.shape)
       out = self.conv2(out)
       out = self.bn2(out)
       #print("out2.shape:", out.shape)
       out += identity
       out = self.relu(out)
       #print("out.shape:", out.shape)
       #print("")
       return out
```

Next we implement a stack of residual blocks for convenience. The first layer in the block is the one changing the number of channels and downsampling. You can use nn.ModuleList to use a list of child modules.

```
[0]: class ResidualStack(nn.Module):
    """
    A stack of residual blocks.

Args:
    in_channels: The number of channels (feature maps) of the incoming
    →embedding
    out_channels: The number of channels after the first layer
    stride: Stride size of the first layer, used for downsampling
```

```
num_blocks: Number of residual blocks
"""

def __init__(self, in_channels, out_channels, stride, num_blocks):
    super().__init__()

# TODO: Initialize the required layers (blocks)
    self.layers = nn.ModuleList([ResidualBlock(in_channels, out_channels,])

for i in range(1, num_blocks):
    self.layers.append(ResidualBlock(out_channels, out_channels, 1))

def forward(self, input):
    # TODO: Execute the layers (blocks)
    #print("Execute " + str(len(self.layers)) + " ResidualBlock layers")
    for layer in self.layers:
        input = layer(input)
    return input
```

Now we are finally ready to implement the full model! To do this, use the nn.Sequential API and carefully read the following paragraph from the paper (Fig. 3 is not important):

attachment:resnet_cifar10_description.png

Note that a convolution layer is always convolution + batch norm + activation (ReLU), that each ResidualBlock contains 2 layers, and that you might have to squeeze the embedding before the dense (fully-connected) layer.

```
[0]: n = 5
    num_classes = 10
     class Squeeze(torch.nn.Module):
         def forward(self, input):
           out = input.squeeze()
           return out
     # TODO: Implement ResNet via nn.Sequential
     resnet = nn.Sequential(
                 ResidualStack(3, 16, 1, 1),
                 ResidualStack(16, 16, 1, 2 * n),
                 ResidualStack(16, 32, 2, 2 * n),
                 ResidualStack(32, 64, 2, 2 * n),
                 nn.AdaptiveAvgPool2d(1),
                 #nn.AdaptiveAvgPool2d(64),
                 Squeeze(),
                 nn.Linear(64, num_classes)
```

Next we need to initialize the weights of our model.

```
[0]: def initialize_weight(module):
    if isinstance(module, (nn.Linear, nn.Conv2d)):
        nn.init.kaiming_normal_(module.weight, nonlinearity='relu')
    elif isinstance(module, nn.BatchNorm2d):
        nn.init.constant_(module.weight, 1)
        nn.init.constant_(module.bias, 0)

resnet.apply(initialize_weight);
```

5 4. Training

So now we have a shiny new model, but that doesn't really help when we can't train it. So that's what we do next.

First we need to load the data. Note that we split the official training data into train and validation sets, because you must not look at the test set until you are completely done developing your model and report the final results. Some people don't do this properly, but you should not copy other people's bad habits.

```
class CIFAR10Subset(torchvision.datasets.CIFAR10):
    """
    Get a subset of the CIFAR10 dataset, according to the passed indices.
    """
    def __init__(self, *args, idx=None, **kwargs):
        super().__init__(*args, **kwargs)

    if idx is None:
        return

    self.data = self.data[idx]
    targets_np = np.array(self.targets)
    self.targets = targets_np[idx].tolist()
```

We next define transformations that change the images into PyTorch tensors, standardize the values according to the precomputed mean and standard deviation, and provide data augmentation for the training set.

```
normalize
      ])
[14]: ntrain = 45 000
      train_set = CIFAR10Subset(root='./data', train=True, idx=range(ntrain),
                                download=True, transform=transform_train)
      val_set = CIFAR10Subset(root='./data', train=True, idx=range(ntrain, 50_000),
                              download=True, transform=transform_eval)
      test_set = torchvision.datasets.CIFAR10(root='./data', train=False,
                                              download=True, transform=transform_eval)
     0it [00:00, ?it/s]
     Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to
     ./data/cifar-10-python.tar.gz
     170500096it [00:06, 27285630.82it/s]
     Extracting ./data/cifar-10-python.tar.gz to ./data
     Files already downloaded and verified
     Files already downloaded and verified
 [0]: dataloaders = {}
      dataloaders['train'] = torch.utils.data.DataLoader(train set, batch size=128,
                                                         shuffle=True, num_workers=0,
                                                         pin memory=True)
      dataloaders['val'] = torch.utils.data.DataLoader(val_set, batch_size=128,
                                                       shuffle=False, num_workers=0,
                                                       pin_memory=True)
```

Next we push the model to our GPU (if there is one).

dataloaders['test'] = torch.utils.data.DataLoader(test_set, batch_size=128,

shuffle=False, num_workers=0,

pin_memory=True)

True

Next we define a helper method that does one epoch of training or evaluation. We have only defined training here, so you need to implement the necessary changes for evaluation!

```
[0]: def run_epoch(model, optimizer, dataloader, train):
    """
    Run one epoch of training or evaluation.

Args:
```

```
model: The model used for prediction
       optimizer: Optimization algorithm for the model
       dataloader: Dataloader providing the data to run our model on
       train: Whether this epoch is used for training or evaluation
  Returns:
      Loss and accuracy in this epoch.
   # TODO: Change the necessary parts to work correctly during evaluation
\hookrightarrow (train=False)
  device = next(model.parameters()).device
  # Set model to training mode (for e.g. batch normalization, dropout)
  model.train()
  epoch_loss = 0.0
  epoch_acc = 0.0
   # Iterate over data
  for xb, yb in dataloader:
       xb, yb = xb.to(device), yb.to(device)
       # zero the parameter gradients
       if train != False:
           optimizer.zero_grad()
       # forward
       with torch.set_grad_enabled(True):
           pred = model(xb)
           loss = F.cross_entropy(pred, yb)
           top1 = torch.argmax(pred, dim=1)
           ncorrect = torch.sum(top1 == yb)
           if train != False:
               loss.backward()
               optimizer.step()
       # statistics
       epoch_loss += loss.item()
       epoch_acc += ncorrect.item()
   epoch_loss /= len(dataloader.dataset)
   epoch_acc /= len(dataloader.dataset)
  return epoch_loss, epoch_acc
```

Next we implement a method for fitting (training) our model. For many models early stopping

can save a lot of training time. Your task is to add early stopping to the loop (based on validation accuracy)! And don't forget to save the best model parameters according to validation accuracy. You will need copy.deepcopy and the state_dict for this.

```
[0]: def fit(model, optimizer, lr_scheduler, dataloaders, max_epochs, patience):
         Fit the given model on the dataset.
         Args:
             model: The model used for prediction
             optimizer: Optimization algorithm for the model
             lr_scheduler: Learning rate scheduler that improves training
                            in late epochs with learning rate decay
             dataloaders: Dataloaders for training and validation
             max_epochs: Maximum number of epochs for training
             patience: Number of epochs to wait with early stopping the
                        training if validation loss has decreased
         Returns:
             Loss and accuracy in this epoch.
         best_acc = 0.0
         curr patience = 0
         best_model_weights = None
         for epoch in range(max_epochs):
             train_loss, train_acc = run_epoch(model, optimizer,_

dataloaders['train'], train=True)

             lr_scheduler.step()
             print(f"Epoch {epoch + 1: >3}/{max_epochs}, train loss: {train_loss:.
      \rightarrow2e}, accuracy: {train_acc * 100:.2f}%")
             val_loss, val_acc = run_epoch(model, None, dataloaders['val'],__
      →train=False)
             print(f"Epoch {epoch + 1: >3}/{max_epochs}, val loss: {val_loss:.2e},__
      \rightarrowaccuracy: {val acc * 100:.2f}%")
             # TODO: Add early stopping and save the best weights (in_
      \rightarrow best_model_weights)
             curr_patience = 1 if val_acc < best_acc else curr_patience + 1</pre>
             updated model weights = False
             if best_acc < val_acc:</pre>
                 best model weights = copy.deepcopy(model.state dict())
                 updated model weights = True
             best_acc = max(val_acc, best_acc)
```

```
print(f"Updates: curr_patience={curr_patience}, best_acc={best_acc},

updated={updated_model_weights}")

if curr_patience > patience or epoch + 1 >= max_epochs:
    print(f"Stop early: curr_patience > patience or epoch + 1 >= u

max_epochs")
    break

model.load_state_dict(best_model_weights)
```

In most cases you should just use the Adam optimizer for training, because it works well out of the box. However, a well-tuned SGD (with momentum) will in most cases outperform Adam. And since the original paper gives us a well-tuned SGD we will just use that.

```
1/200, train loss: 1.80e-02, accuracy: 18.42%
Epoch
        1/200, val loss: 1.49e-02, accuracy: 29.16%
Epoch
Updates: curr_patience=1, best_acc=0.2916, updated=True
       2/200, train loss: 1.33e-02, accuracy: 36.39%
Epoch
       2/200, val loss: 1.22e-02, accuracy: 43.74%
Epoch
Updates: curr_patience=2, best_acc=0.4374, updated=True
Epoch
       3/200, train loss: 1.10e-02, accuracy: 48.72%
Epoch
       3/200, val loss: 1.02e-02, accuracy: 56.32%
Updates: curr_patience=3, best_acc=0.5632, updated=True
Epoch
       4/200, train loss: 9.33e-03, accuracy: 57.26%
Epoch
       4/200, val loss: 8.92e-03, accuracy: 61.78%
Updates: curr patience=4, best acc=0.6178, updated=True
       5/200, train loss: 8.20e-03, accuracy: 62.74%
Epoch
        5/200, val loss: 8.05e-03, accuracy: 66.16%
Epoch
Updates: curr_patience=5, best_acc=0.6616, updated=True
Epoch
        6/200, train loss: 7.22e-03, accuracy: 67.73%
Epoch
        6/200, val loss: 7.06e-03, accuracy: 70.32%
Updates: curr_patience=6, best_acc=0.7032, updated=True
       7/200, train loss: 6.42e-03, accuracy: 71.36%
Epoch
Epoch
       7/200, val loss: 6.53e-03, accuracy: 73.50%
Updates: curr_patience=7, best_acc=0.735, updated=True
Epoch
       8/200, train loss: 5.82e-03, accuracy: 74.03%
Epoch
       8/200, val loss: 6.09e-03, accuracy: 74.92%
Updates: curr_patience=8, best_acc=0.7492, updated=True
       9/200, train loss: 5.37e-03, accuracy: 76.19%
Epoch
```

9/200, val loss: 5.90e-03, accuracy: 76.60% Updates: curr_patience=9, best_acc=0.766, updated=True Epoch 10/200, train loss: 4.98e-03, accuracy: 77.90% Epoch 10/200, val loss: 5.14e-03, accuracy: 78.54% Updates: curr patience=10, best acc=0.7854, updated=True Epoch 11/200, train loss: 4.67e-03, accuracy: 79.30% Epoch 11/200, val loss: 4.94e-03, accuracy: 79.10% Updates: curr_patience=11, best_acc=0.791, updated=True Epoch 12/200, train loss: 4.41e-03, accuracy: 80.37% Epoch 12/200, val loss: 5.12e-03, accuracy: 79.14% Updates: curr_patience=12, best_acc=0.7914, updated=True Epoch 13/200, train loss: 4.18e-03, accuracy: 81.44% Epoch 13/200, val loss: 4.56e-03, accuracy: 80.24% Updates: curr_patience=13, best_acc=0.8024, updated=True Epoch 14/200, train loss: 3.97e-03, accuracy: 82.44% Epoch 14/200, val loss: 4.46e-03, accuracy: 81.68% Updates: curr_patience=14, best_acc=0.8168, updated=True Epoch 15/200, train loss: 3.85e-03, accuracy: 82.88% Epoch 15/200, val loss: 4.50e-03, accuracy: 81.58% Updates: curr patience=1, best acc=0.8168, updated=False Epoch 16/200, train loss: 3.64e-03, accuracy: 83.89% Epoch 16/200, val loss: 4.29e-03, accuracy: 82.32% Updates: curr_patience=2, best_acc=0.8232, updated=True Epoch 17/200, train loss: 3.46e-03, accuracy: 84.62% Epoch 17/200, val loss: 4.17e-03, accuracy: 83.06% Updates: curr_patience=3, best_acc=0.8306, updated=True Epoch 18/200, train loss: 3.39e-03, accuracy: 85.13% Epoch 18/200, val loss: 4.26e-03, accuracy: 82.46% Updates: curr_patience=1, best_acc=0.8306, updated=False Epoch 19/200, train loss: 3.26e-03, accuracy: 85.57% Epoch 19/200, val loss: 4.17e-03, accuracy: 83.46% Updates: curr_patience=2, best_acc=0.8346, updated=True Epoch 20/200, train loss: 3.14e-03, accuracy: 86.09% Epoch 20/200, val loss: 4.19e-03, accuracy: 83.56% Updates: curr patience=3, best acc=0.8356, updated=True Epoch 21/200, train loss: 3.08e-03, accuracy: 86.30% Epoch 21/200, val loss: 4.05e-03, accuracy: 83.22% Updates: curr_patience=1, best_acc=0.8356, updated=False Epoch 22/200, train loss: 2.98e-03, accuracy: 86.71% Epoch 22/200, val loss: 3.97e-03, accuracy: 83.44% Updates: curr_patience=1, best_acc=0.8356, updated=False Epoch 23/200, train loss: 2.91e-03, accuracy: 87.10% Epoch 23/200, val loss: 4.06e-03, accuracy: 83.42% Updates: curr_patience=1, best_acc=0.8356, updated=False Epoch 24/200, train loss: 2.83e-03, accuracy: 87.44% Epoch 24/200, val loss: 3.85e-03, accuracy: 83.90% Updates: curr_patience=2, best_acc=0.839, updated=True Epoch 25/200, train loss: 2.79e-03, accuracy: 87.55%

Epoch 25/200, val loss: 3.73e-03, accuracy: 84.46% Updates: curr_patience=3, best_acc=0.8446, updated=True Epoch 26/200, train loss: 2.76e-03, accuracy: 87.70% Epoch 26/200, val loss: 3.97e-03, accuracy: 83.82% Updates: curr patience=1, best acc=0.8446, updated=False Epoch 27/200, train loss: 2.66e-03, accuracy: 88.26% Epoch 27/200, val loss: 3.75e-03, accuracy: 84.78% Updates: curr_patience=2, best_acc=0.8478, updated=True Epoch 28/200, train loss: 2.59e-03, accuracy: 88.55% Epoch 28/200, val loss: 3.61e-03, accuracy: 84.90% Updates: curr_patience=3, best_acc=0.849, updated=True Epoch 29/200, train loss: 2.52e-03, accuracy: 88.80% Epoch 29/200, val loss: 3.62e-03, accuracy: 85.06% Updates: curr_patience=4, best_acc=0.8506, updated=True Epoch 30/200, train loss: 2.50e-03, accuracy: 88.82% Epoch 30/200, val loss: 3.58e-03, accuracy: 85.20% Updates: curr_patience=5, best_acc=0.852, updated=True Epoch 31/200, train loss: 2.47e-03, accuracy: 89.05% Epoch 31/200, val loss: 3.59e-03, accuracy: 84.94% Updates: curr patience=1, best acc=0.852, updated=False Epoch 32/200, train loss: 2.41e-03, accuracy: 89.30% Epoch 32/200, val loss: 3.64e-03, accuracy: 85.38% Updates: curr_patience=2, best_acc=0.8538, updated=True Epoch 33/200, train loss: 2.35e-03, accuracy: 89.53% Epoch 33/200, val loss: 3.55e-03, accuracy: 86.08% Updates: curr_patience=3, best_acc=0.8608, updated=True Epoch 34/200, train loss: 2.34e-03, accuracy: 89.71% Epoch 34/200, val loss: 3.59e-03, accuracy: 85.30% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 35/200, train loss: 2.31e-03, accuracy: 89.72% Epoch 35/200, val loss: 3.56e-03, accuracy: 85.54% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 36/200, train loss: 2.25e-03, accuracy: 89.81% Epoch 36/200, val loss: 3.66e-03, accuracy: 85.32% Updates: curr patience=1, best acc=0.8608, updated=False Epoch 37/200, train loss: 2.22e-03, accuracy: 89.99% Epoch 37/200, val loss: 3.77e-03, accuracy: 84.86% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 38/200, train loss: 2.26e-03, accuracy: 89.89% Epoch 38/200, val loss: 3.47e-03, accuracy: 85.88% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 39/200, train loss: 2.13e-03, accuracy: 90.43% Epoch 39/200, val loss: 3.66e-03, accuracy: 85.62% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 40/200, train loss: 2.13e-03, accuracy: 90.50% Epoch 40/200, val loss: 3.66e-03, accuracy: 85.08% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 41/200, train loss: 2.10e-03, accuracy: 90.51%

Epoch 41/200, val loss: 3.49e-03, accuracy: 85.82% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 42/200, train loss: 2.07e-03, accuracy: 90.70% Epoch 42/200, val loss: 3.41e-03, accuracy: 85.86% Updates: curr patience=1, best acc=0.8608, updated=False Epoch 43/200, train loss: 2.04e-03, accuracy: 91.04% Epoch 43/200, val loss: 3.64e-03, accuracy: 85.92% Updates: curr_patience=1, best_acc=0.8608, updated=False Epoch 44/200, train loss: 2.02e-03, accuracy: 90.98% Epoch 44/200, val loss: 3.52e-03, accuracy: 86.28% Updates: curr_patience=2, best_acc=0.8628, updated=True Epoch 45/200, train loss: 1.97e-03, accuracy: 91.14% Epoch 45/200, val loss: 3.73e-03, accuracy: 85.34% Updates: curr_patience=1, best_acc=0.8628, updated=False Epoch 46/200, train loss: 1.95e-03, accuracy: 91.30% Epoch 46/200, val loss: 3.82e-03, accuracy: 85.20% Updates: curr_patience=1, best_acc=0.8628, updated=False Epoch 47/200, train loss: 1.98e-03, accuracy: 91.17% Epoch 47/200, val loss: 3.38e-03, accuracy: 86.98% Updates: curr patience=2, best acc=0.8698, updated=True Epoch 48/200, train loss: 1.93e-03, accuracy: 91.34% Epoch 48/200, val loss: 3.77e-03, accuracy: 85.46% Updates: curr_patience=1, best_acc=0.8698, updated=False Epoch 49/200, train loss: 1.91e-03, accuracy: 91.52% Epoch 49/200, val loss: 3.58e-03, accuracy: 85.54% Updates: curr_patience=1, best_acc=0.8698, updated=False Epoch 50/200, train loss: 1.87e-03, accuracy: 91.49% Epoch 50/200, val loss: 3.35e-03, accuracy: 86.56% Updates: curr_patience=1, best_acc=0.8698, updated=False Epoch 51/200, train loss: 1.86e-03, accuracy: 91.70% Epoch 51/200, val loss: 3.23e-03, accuracy: 87.06% Updates: curr_patience=2, best_acc=0.8706, updated=True Epoch 52/200, train loss: 1.89e-03, accuracy: 91.59% Epoch 52/200, val loss: 3.36e-03, accuracy: 87.02% Updates: curr patience=1, best acc=0.8706, updated=False Epoch 53/200, train loss: 1.82e-03, accuracy: 91.88% Epoch 53/200, val loss: 3.52e-03, accuracy: 86.30% Updates: curr_patience=1, best_acc=0.8706, updated=False Epoch 54/200, train loss: 1.81e-03, accuracy: 92.00% Epoch 54/200, val loss: 3.50e-03, accuracy: 86.86% Updates: curr_patience=1, best_acc=0.8706, updated=False Epoch 55/200, train loss: 1.81e-03, accuracy: 91.93% Epoch 55/200, val loss: 3.45e-03, accuracy: 86.64% Updates: curr_patience=1, best_acc=0.8706, updated=False Epoch 56/200, train loss: 1.80e-03, accuracy: 91.83% Epoch 56/200, val loss: 3.57e-03, accuracy: 86.52% Updates: curr_patience=1, best_acc=0.8706, updated=False Epoch 57/200, train loss: 1.72e-03, accuracy: 92.24%

Epoch 57/200, val loss: 3.13e-03, accuracy: 87.72% Updates: curr_patience=2, best_acc=0.8772, updated=True Epoch 58/200, train loss: 1.74e-03, accuracy: 92.21% Epoch 58/200, val loss: 3.41e-03, accuracy: 86.68% Updates: curr patience=1, best acc=0.8772, updated=False Epoch 59/200, train loss: 1.74e-03, accuracy: 92.22% Epoch 59/200, val loss: 3.19e-03, accuracy: 87.32% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 60/200, train loss: 1.69e-03, accuracy: 92.51% Epoch 60/200, val loss: 3.34e-03, accuracy: 86.86% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 61/200, train loss: 1.71e-03, accuracy: 92.39% Epoch 61/200, val loss: 3.46e-03, accuracy: 87.16% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 62/200, train loss: 1.71e-03, accuracy: 92.21% Epoch 62/200, val loss: 3.41e-03, accuracy: 86.38% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 63/200, train loss: 1.71e-03, accuracy: 92.17% Epoch 63/200, val loss: 3.24e-03, accuracy: 86.92% Updates: curr patience=1, best acc=0.8772, updated=False Epoch 64/200, train loss: 1.67e-03, accuracy: 92.52% Epoch 64/200, val loss: 3.53e-03, accuracy: 86.58% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 65/200, train loss: 1.65e-03, accuracy: 92.71% Epoch 65/200, val loss: 3.40e-03, accuracy: 87.32% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 66/200, train loss: 1.67e-03, accuracy: 92.53% Epoch 66/200, val loss: 3.47e-03, accuracy: 87.02% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 67/200, train loss: 1.59e-03, accuracy: 92.79% Epoch 67/200, val loss: 3.63e-03, accuracy: 86.06% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 68/200, train loss: 1.69e-03, accuracy: 92.32% Epoch 68/200, val loss: 3.37e-03, accuracy: 87.30% Updates: curr patience=1, best acc=0.8772, updated=False Epoch 69/200, train loss: 1.58e-03, accuracy: 92.79% Epoch 69/200, val loss: 3.41e-03, accuracy: 86.68% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 70/200, train loss: 1.61e-03, accuracy: 92.82% Epoch 70/200, val loss: 3.25e-03, accuracy: 87.28% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 71/200, train loss: 1.57e-03, accuracy: 92.92% Epoch 71/200, val loss: 3.47e-03, accuracy: 86.84% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 72/200, train loss: 1.61e-03, accuracy: 92.76% Epoch 72/200, val loss: 3.41e-03, accuracy: 86.86% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 73/200, train loss: 1.57e-03, accuracy: 92.85%

Epoch 73/200, val loss: 3.42e-03, accuracy: 87.22% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 74/200, train loss: 1.56e-03, accuracy: 92.99% Epoch 74/200, val loss: 3.34e-03, accuracy: 87.14% Updates: curr patience=1, best acc=0.8772, updated=False Epoch 75/200, train loss: 1.52e-03, accuracy: 93.29% Epoch 75/200, val loss: 3.32e-03, accuracy: 87.04% Updates: curr_patience=1, best_acc=0.8772, updated=False Epoch 76/200, train loss: 1.48e-03, accuracy: 93.42% Epoch 76/200, val loss: 3.25e-03, accuracy: 88.00% Updates: curr_patience=2, best_acc=0.88, updated=True Epoch 77/200, train loss: 1.56e-03, accuracy: 93.10% Epoch 77/200, val loss: 3.36e-03, accuracy: 86.92% Updates: curr_patience=1, best_acc=0.88, updated=False Epoch 78/200, train loss: 1.52e-03, accuracy: 93.17% Epoch 78/200, val loss: 3.47e-03, accuracy: 86.96% Updates: curr_patience=1, best_acc=0.88, updated=False Epoch 79/200, train loss: 1.51e-03, accuracy: 93.28% Epoch 79/200, val loss: 3.39e-03, accuracy: 86.76% Updates: curr patience=1, best acc=0.88, updated=False Epoch 80/200, train loss: 1.51e-03, accuracy: 93.23% Epoch 80/200, val loss: 2.96e-03, accuracy: 88.20% Updates: curr_patience=2, best_acc=0.882, updated=True Epoch 81/200, train loss: 1.51e-03, accuracy: 93.16% Epoch 81/200, val loss: 3.22e-03, accuracy: 87.18% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 82/200, train loss: 1.47e-03, accuracy: 93.48% Epoch 82/200, val loss: 3.16e-03, accuracy: 88.14% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 83/200, train loss: 1.45e-03, accuracy: 93.40% Epoch 83/200, val loss: 3.43e-03, accuracy: 86.44% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 84/200, train loss: 1.51e-03, accuracy: 93.24% Epoch 84/200, val loss: 3.21e-03, accuracy: 87.50% Updates: curr patience=1, best acc=0.882, updated=False Epoch 85/200, train loss: 1.47e-03, accuracy: 93.30% Epoch 85/200, val loss: 3.37e-03, accuracy: 87.58% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 86/200, train loss: 1.50e-03, accuracy: 93.12% Epoch 86/200, val loss: 3.38e-03, accuracy: 87.70% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 87/200, train loss: 1.49e-03, accuracy: 93.28% Epoch 87/200, val loss: 3.39e-03, accuracy: 87.30% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 88/200, train loss: 1.43e-03, accuracy: 93.54% Epoch 88/200, val loss: 3.27e-03, accuracy: 87.44% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 89/200, train loss: 1.50e-03, accuracy: 93.20%

Epoch 89/200, val loss: 3.42e-03, accuracy: 87.02% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 90/200, train loss: 1.48e-03, accuracy: 93.43% Epoch 90/200, val loss: 3.18e-03, accuracy: 88.12% Updates: curr patience=1, best acc=0.882, updated=False Epoch 91/200, train loss: 1.47e-03, accuracy: 93.41% Epoch 91/200, val loss: 3.07e-03, accuracy: 88.04% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 92/200, train loss: 1.43e-03, accuracy: 93.60% Epoch 92/200, val loss: 3.27e-03, accuracy: 87.80% Updates: curr_patience=1, best_acc=0.882, updated=False Epoch 93/200, train loss: 1.45e-03, accuracy: 93.43% Epoch 93/200, val loss: 3.30e-03, accuracy: 88.28% Updates: curr_patience=2, best_acc=0.8828, updated=True Epoch 94/200, train loss: 1.43e-03, accuracy: 93.68% Epoch 94/200, val loss: 3.32e-03, accuracy: 87.36% Updates: curr_patience=1, best_acc=0.8828, updated=False Epoch 95/200, train loss: 1.42e-03, accuracy: 93.63% Epoch 95/200, val loss: 3.28e-03, accuracy: 88.18% Updates: curr patience=1, best acc=0.8828, updated=False Epoch 96/200, train loss: 1.43e-03, accuracy: 93.50% Epoch 96/200, val loss: 3.59e-03, accuracy: 87.06% Updates: curr_patience=1, best_acc=0.8828, updated=False Epoch 97/200, train loss: 1.43e-03, accuracy: 93.69% Epoch 97/200, val loss: 3.36e-03, accuracy: 87.60% Updates: curr_patience=1, best_acc=0.8828, updated=False Epoch 98/200, train loss: 1.45e-03, accuracy: 93.45% Epoch 98/200, val loss: 3.55e-03, accuracy: 86.86% Updates: curr_patience=1, best_acc=0.8828, updated=False Epoch 99/200, train loss: 1.39e-03, accuracy: 93.53% Epoch 99/200, val loss: 3.01e-03, accuracy: 88.44% Updates: curr_patience=2, best_acc=0.8844, updated=True Epoch 100/200, train loss: 1.38e-03, accuracy: 93.80% Epoch 100/200, val loss: 3.61e-03, accuracy: 86.70% Updates: curr patience=1, best acc=0.8844, updated=False Epoch 101/200, train loss: 7.20e-04, accuracy: 96.86% Epoch 101/200, val loss: 2.71e-03, accuracy: 89.98% Updates: curr_patience=2, best_acc=0.8998, updated=True Epoch 102/200, train loss: 4.79e-04, accuracy: 98.01% Epoch 102/200, val loss: 2.67e-03, accuracy: 90.60% Updates: curr_patience=3, best_acc=0.906, updated=True Epoch 103/200, train loss: 4.08e-04, accuracy: 98.33% Epoch 103/200, val loss: 2.69e-03, accuracy: 90.68% Updates: curr_patience=4, best_acc=0.9068, updated=True Epoch 104/200, train loss: 3.61e-04, accuracy: 98.54% Epoch 104/200, val loss: 2.77e-03, accuracy: 90.70% Updates: curr_patience=5, best_acc=0.907, updated=True Epoch 105/200, train loss: 3.31e-04, accuracy: 98.62%

Epoch 105/200, val loss: 2.77e-03, accuracy: 90.70% Updates: curr_patience=6, best_acc=0.907, updated=False Epoch 106/200, train loss: 2.93e-04, accuracy: 98.82% Epoch 106/200, val loss: 2.89e-03, accuracy: 90.58% Updates: curr patience=1, best acc=0.907, updated=False Epoch 107/200, train loss: 2.73e-04, accuracy: 98.89% Epoch 107/200, val loss: 2.88e-03, accuracy: 90.98% Updates: curr_patience=2, best_acc=0.9098, updated=True Epoch 108/200, train loss: 2.46e-04, accuracy: 98.98% Epoch 108/200, val loss: 2.92e-03, accuracy: 90.92% Updates: curr_patience=1, best_acc=0.9098, updated=False Epoch 109/200, train loss: 2.32e-04, accuracy: 99.04% Epoch 109/200, val loss: 2.95e-03, accuracy: 90.86% Updates: curr_patience=1, best_acc=0.9098, updated=False Epoch 110/200, train loss: 1.99e-04, accuracy: 99.18% Epoch 110/200, val loss: 3.06e-03, accuracy: 90.78% Updates: curr_patience=1, best_acc=0.9098, updated=False Epoch 111/200, train loss: 1.95e-04, accuracy: 99.27% Epoch 111/200, val loss: 3.08e-03, accuracy: 90.64% Updates: curr patience=1, best acc=0.9098, updated=False Epoch 112/200, train loss: 1.89e-04, accuracy: 99.27% Epoch 112/200, val loss: 3.04e-03, accuracy: 91.00% Updates: curr_patience=2, best_acc=0.91, updated=True Epoch 113/200, train loss: 1.70e-04, accuracy: 99.38% Epoch 113/200, val loss: 3.11e-03, accuracy: 90.62% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 114/200, train loss: 1.74e-04, accuracy: 99.31% Epoch 114/200, val loss: 3.11e-03, accuracy: 90.70% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 115/200, train loss: 1.69e-04, accuracy: 99.33% Epoch 115/200, val loss: 3.19e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 116/200, train loss: 1.52e-04, accuracy: 99.42% Epoch 116/200, val loss: 3.29e-03, accuracy: 90.48% Updates: curr patience=1, best acc=0.91, updated=False Epoch 117/200, train loss: 1.43e-04, accuracy: 99.44% Epoch 117/200, val loss: 3.37e-03, accuracy: 90.54% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 118/200, train loss: 1.37e-04, accuracy: 99.48% Epoch 118/200, val loss: 3.31e-03, accuracy: 90.36% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 119/200, train loss: 1.31e-04, accuracy: 99.48% Epoch 119/200, val loss: 3.29e-03, accuracy: 90.68% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 120/200, train loss: 1.27e-04, accuracy: 99.50% Epoch 120/200, val loss: 3.30e-03, accuracy: 90.62% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 121/200, train loss: 1.31e-04, accuracy: 99.47%

Epoch 121/200, val loss: 3.48e-03, accuracy: 90.50% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 122/200, train loss: 1.12e-04, accuracy: 99.60% Epoch 122/200, val loss: 3.40e-03, accuracy: 90.68% Updates: curr patience=1, best acc=0.91, updated=False Epoch 123/200, train loss: 1.15e-04, accuracy: 99.52% Epoch 123/200, val loss: 3.55e-03, accuracy: 90.34% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 124/200, train loss: 1.09e-04, accuracy: 99.59% Epoch 124/200, val loss: 3.42e-03, accuracy: 90.62% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 125/200, train loss: 1.06e-04, accuracy: 99.59% Epoch 125/200, val loss: 3.45e-03, accuracy: 90.58% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 126/200, train loss: 9.50e-05, accuracy: 99.66% Epoch 126/200, val loss: 3.41e-03, accuracy: 90.88% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 127/200, train loss: 1.05e-04, accuracy: 99.64% Epoch 127/200, val loss: 3.53e-03, accuracy: 90.42% Updates: curr patience=1, best acc=0.91, updated=False Epoch 128/200, train loss: 9.98e-05, accuracy: 99.62% Epoch 128/200, val loss: 3.42e-03, accuracy: 90.54% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 129/200, train loss: 1.00e-04, accuracy: 99.58% Epoch 129/200, val loss: 3.55e-03, accuracy: 90.66% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 130/200, train loss: 8.87e-05, accuracy: 99.68% Epoch 130/200, val loss: 3.49e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 131/200, train loss: 8.15e-05, accuracy: 99.73% Epoch 131/200, val loss: 3.52e-03, accuracy: 90.76% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 132/200, train loss: 8.20e-05, accuracy: 99.70% Epoch 132/200, val loss: 3.52e-03, accuracy: 90.74% Updates: curr patience=1, best acc=0.91, updated=False Epoch 133/200, train loss: 7.63e-05, accuracy: 99.74% Epoch 133/200, val loss: 3.59e-03, accuracy: 90.70% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 134/200, train loss: 8.69e-05, accuracy: 99.69% Epoch 134/200, val loss: 3.54e-03, accuracy: 90.84% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 135/200, train loss: 7.38e-05, accuracy: 99.76% Epoch 135/200, val loss: 3.54e-03, accuracy: 90.74% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 136/200, train loss: 7.20e-05, accuracy: 99.72% Epoch 136/200, val loss: 3.64e-03, accuracy: 90.50% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 137/200, train loss: 7.26e-05, accuracy: 99.76%

Epoch 137/200, val loss: 3.79e-03, accuracy: 90.32% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 138/200, train loss: 7.19e-05, accuracy: 99.76% Epoch 138/200, val loss: 3.56e-03, accuracy: 90.62% Updates: curr patience=1, best acc=0.91, updated=False Epoch 139/200, train loss: 7.23e-05, accuracy: 99.74% Epoch 139/200, val loss: 3.60e-03, accuracy: 90.44% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 140/200, train loss: 7.39e-05, accuracy: 99.74% Epoch 140/200, val loss: 3.68e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.91, updated=False Epoch 141/200, train loss: 7.02e-05, accuracy: 99.77% Epoch 141/200, val loss: 3.74e-03, accuracy: 91.10% Updates: curr_patience=2, best_acc=0.911, updated=True Epoch 142/200, train loss: 6.90e-05, accuracy: 99.75% Epoch 142/200, val loss: 3.73e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 143/200, train loss: 6.73e-05, accuracy: 99.77% Epoch 143/200, val loss: 3.74e-03, accuracy: 90.86% Updates: curr patience=1, best acc=0.911, updated=False Epoch 144/200, train loss: 6.83e-05, accuracy: 99.78% Epoch 144/200, val loss: 3.70e-03, accuracy: 90.84% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 145/200, train loss: 6.36e-05, accuracy: 99.78% Epoch 145/200, val loss: 3.64e-03, accuracy: 91.04% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 146/200, train loss: 6.35e-05, accuracy: 99.78% Epoch 146/200, val loss: 3.82e-03, accuracy: 90.70% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 147/200, train loss: 6.96e-05, accuracy: 99.74% Epoch 147/200, val loss: 3.80e-03, accuracy: 90.50% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 148/200, train loss: 6.83e-05, accuracy: 99.74% Epoch 148/200, val loss: 3.73e-03, accuracy: 90.80% Updates: curr patience=1, best acc=0.911, updated=False Epoch 149/200, train loss: 5.86e-05, accuracy: 99.80% Epoch 149/200, val loss: 3.81e-03, accuracy: 90.68% Updates: curr_patience=1, best_acc=0.911, updated=False Epoch 150/200, train loss: 5.65e-05, accuracy: 99.79% Epoch 150/200, val loss: 3.82e-03, accuracy: 91.18% Updates: curr_patience=2, best_acc=0.9118, updated=True Epoch 151/200, train loss: 5.20e-05, accuracy: 99.80% Epoch 151/200, val loss: 3.79e-03, accuracy: 91.14% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 152/200, train loss: 4.47e-05, accuracy: 99.86% Epoch 152/200, val loss: 3.79e-03, accuracy: 91.16% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 153/200, train loss: 4.52e-05, accuracy: 99.86%

Epoch 153/200, val loss: 3.78e-03, accuracy: 91.16% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 154/200, train loss: 4.77e-05, accuracy: 99.84% Epoch 154/200, val loss: 3.80e-03, accuracy: 90.94% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 155/200, train loss: 4.24e-05, accuracy: 99.87% Epoch 155/200, val loss: 3.79e-03, accuracy: 90.98% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 156/200, train loss: 4.05e-05, accuracy: 99.87% Epoch 156/200, val loss: 3.79e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 157/200, train loss: 3.85e-05, accuracy: 99.91% Epoch 157/200, val loss: 3.79e-03, accuracy: 90.92% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 158/200, train loss: 3.74e-05, accuracy: 99.88% Epoch 158/200, val loss: 3.78e-03, accuracy: 90.88% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 159/200, train loss: 3.89e-05, accuracy: 99.89% Epoch 159/200, val loss: 3.78e-03, accuracy: 90.92% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 160/200, train loss: 3.73e-05, accuracy: 99.91% Epoch 160/200, val loss: 3.79e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 161/200, train loss: 3.73e-05, accuracy: 99.90% Epoch 161/200, val loss: 3.81e-03, accuracy: 90.86% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 162/200, train loss: 3.79e-05, accuracy: 99.88% Epoch 162/200, val loss: 3.79e-03, accuracy: 90.92% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 163/200, train loss: 3.51e-05, accuracy: 99.92% Epoch 163/200, val loss: 3.79e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 164/200, train loss: 3.81e-05, accuracy: 99.90% Epoch 164/200, val loss: 3.78e-03, accuracy: 91.00% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 165/200, train loss: 3.74e-05, accuracy: 99.89% Epoch 165/200, val loss: 3.77e-03, accuracy: 91.10% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 166/200, train loss: 3.34e-05, accuracy: 99.90% Epoch 166/200, val loss: 3.78e-03, accuracy: 91.04% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 167/200, train loss: 3.06e-05, accuracy: 99.93% Epoch 167/200, val loss: 3.80e-03, accuracy: 91.10% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 168/200, train loss: 3.33e-05, accuracy: 99.91% Epoch 168/200, val loss: 3.80e-03, accuracy: 91.04% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 169/200, train loss: 3.57e-05, accuracy: 99.89%

Epoch 169/200, val loss: 3.78e-03, accuracy: 91.14% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 170/200, train loss: 3.24e-05, accuracy: 99.91% Epoch 170/200, val loss: 3.80e-03, accuracy: 91.12% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 171/200, train loss: 3.44e-05, accuracy: 99.91% Epoch 171/200, val loss: 3.80e-03, accuracy: 91.02% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 172/200, train loss: 3.54e-05, accuracy: 99.90% Epoch 172/200, val loss: 3.80e-03, accuracy: 91.08% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 173/200, train loss: 3.16e-05, accuracy: 99.92% Epoch 173/200, val loss: 3.81e-03, accuracy: 91.00% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 174/200, train loss: 3.56e-05, accuracy: 99.90% Epoch 174/200, val loss: 3.80e-03, accuracy: 90.88% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 175/200, train loss: 3.16e-05, accuracy: 99.92% Epoch 175/200, val loss: 3.81e-03, accuracy: 90.88% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 176/200, train loss: 3.28e-05, accuracy: 99.90% Epoch 176/200, val loss: 3.80e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 177/200, train loss: 3.20e-05, accuracy: 99.91% Epoch 177/200, val loss: 3.81e-03, accuracy: 90.88% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 178/200, train loss: 3.41e-05, accuracy: 99.89% Epoch 178/200, val loss: 3.81e-03, accuracy: 90.96% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 179/200, train loss: 3.32e-05, accuracy: 99.90% Epoch 179/200, val loss: 3.81e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 180/200, train loss: 3.01e-05, accuracy: 99.93% Epoch 180/200, val loss: 3.81e-03, accuracy: 90.94% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 181/200, train loss: 3.18e-05, accuracy: 99.92% Epoch 181/200, val loss: 3.82e-03, accuracy: 90.78% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 182/200, train loss: 2.84e-05, accuracy: 99.93% Epoch 182/200, val loss: 3.82e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 183/200, train loss: 3.19e-05, accuracy: 99.92% Epoch 183/200, val loss: 3.81e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 184/200, train loss: 3.21e-05, accuracy: 99.92% Epoch 184/200, val loss: 3.81e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 185/200, train loss: 3.09e-05, accuracy: 99.90%

Epoch 185/200, val loss: 3.82e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 186/200, train loss: 3.11e-05, accuracy: 99.92% Epoch 186/200, val loss: 3.82e-03, accuracy: 90.84% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 187/200, train loss: 3.06e-05, accuracy: 99.92% Epoch 187/200, val loss: 3.82e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 188/200, train loss: 2.91e-05, accuracy: 99.94% Epoch 188/200, val loss: 3.82e-03, accuracy: 90.86% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 189/200, train loss: 3.08e-05, accuracy: 99.94% Epoch 189/200, val loss: 3.83e-03, accuracy: 90.94% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 190/200, train loss: 3.01e-05, accuracy: 99.91% Epoch 190/200, val loss: 3.81e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 191/200, train loss: 2.96e-05, accuracy: 99.93% Epoch 191/200, val loss: 3.81e-03, accuracy: 90.94% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 192/200, train loss: 3.31e-05, accuracy: 99.91% Epoch 192/200, val loss: 3.82e-03, accuracy: 90.86% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 193/200, train loss: 2.78e-05, accuracy: 99.93% Epoch 193/200, val loss: 3.82e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 194/200, train loss: 2.98e-05, accuracy: 99.92% Epoch 194/200, val loss: 3.83e-03, accuracy: 90.72% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 195/200, train loss: 2.58e-05, accuracy: 99.94% Epoch 195/200, val loss: 3.83e-03, accuracy: 90.76% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 196/200, train loss: 2.89e-05, accuracy: 99.93% Epoch 196/200, val loss: 3.82e-03, accuracy: 90.84% Updates: curr patience=1, best acc=0.9118, updated=False Epoch 197/200, train loss: 2.61e-05, accuracy: 99.94% Epoch 197/200, val loss: 3.83e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 198/200, train loss: 2.73e-05, accuracy: 99.94% Epoch 198/200, val loss: 3.84e-03, accuracy: 90.80% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 199/200, train loss: 2.86e-05, accuracy: 99.93% Epoch 199/200, val loss: 3.83e-03, accuracy: 90.82% Updates: curr_patience=1, best_acc=0.9118, updated=False Epoch 200/200, train loss: 3.48e-05, accuracy: 99.90% Epoch 200/200, val loss: 3.82e-03, accuracy: 90.88% Updates: curr_patience=1, best_acc=0.9118, updated=False Stop early: curr_patience > patience or epoch + 1 >= max_epochs Once the model is trained we run it on the test set to obtain our final accuracy. Note that we can only look at the test set once, everything else would lead to overfitting. So you *must* ignore the test set while developing your model!

```
[20]: test_loss, test_acc = run_epoch(resnet, None, dataloaders['test'], train=False) print(f"Test loss: {test_loss:.1e}, accuracy: {test_acc * 100:.2f}%")
```

Test loss: 3.6e-03, accuracy: 90.36%

That's almost what was reported in the paper (92.49%) and we didn't even train on the full training set.

6 Optional task: Squeeze out all the juice!

Can you do even better? Have a look at A Recipe for Training Neural Networks and at the EfficientNet architecture we discussed in the lecture. Play around with the possibilities PyTorch offers you and see how close you can get to the state of the art on CIFAR-10.

Hint: You can use Google Colab to access some free GPUs for your experiments.