

Machine Learning

Lecture 8 & 9: Deep Learning

Prof. Dr. Stephan Günnemann

Data Mining and Analytics
Technische Universität München

10.12.2018 & 17.12.2018

Reading material

- Goodfellow, Deep Learning: chapters 6, 7, 11
- Bishop: chapters 5.1, 5.2, 5.3, 5.5

Acknowledgements

- Slides based on an earlier version by Patrick van der Smagt


Another look at Logistic Regression

We had before:

$$y \mid \mathbf{x} \sim \text{Bernoulli}\left(\sigma(\mathbf{w}^T \mathbf{x} + w_0)\right)$$
$$\mathbf{w}^T \mathbf{x} := w_0 + w_1 x_1 + \dots + w_D x_D$$

We can represent this graphically
(the first node $1 = x_0$ is for the bias term):

- each node is a (scalar) input
- multiply the input with the weight on the edge: $x_i w_i$
- compute weighted sum of incoming edges:
 $a_0 = \sum_{i=0}^D x_i w_i$
- apply (activation) function: $y = \sigma(a_0)$





The XOR dataset

The XOR dataset is not linearly separable
→ Logistic Regression will fail since it learns a linear decision boundary

In general:

- $\mathcal{X} = \{\mathbf{x}_1, \mathbf{x}_1, \dots, \mathbf{x}_n\} \in \mathbb{R}^d$
- Given n points there are 2^n dichotomies
- Only $2 \cdot \sum_{i=0}^d \binom{n-1}{i}$ are linearly separable¹
- With $n > d$ the probability that \mathcal{X} is linearly separable goes to 0

$O(2^n)$



How to handle non-linearity? Basis functions

We have input vectors \mathbf{x} and associated output values y . We want to describe the underlying functional relation.

We can use the following simple model:

$$f(\mathbf{x}, \mathbf{w}) = \sigma\left(w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})\right) = \sigma(\mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})) \quad (1)$$




where

- $\boldsymbol{\phi}$ **basis function** — many choices, can be nonlinear
- w_0 **bias** — equivalent to defining $\phi_0 \equiv 1$;
or adding constant 1 to every sample

Remember we are **linear** in \mathbf{w} !

Example: Handling XOR dataset by custom basis functions

Apply a (nonlinear) transformation ϕ that maps samples to a space where they are linearly separable. For example:



Here we defined a custom basis function $\phi : \mathbb{R}^3 \rightarrow \mathbb{R}^2$
 $\phi(\mathbf{x}) = \phi(1, x_1, x_2) = (\sigma(5 + x_1 + x_2), \sigma(5 - x_1 - x_2))$

Example: Handling XOR dataset by custom basis functions

Overall function that is modeled:


$$f(\mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \phi(\mathbf{x})) = \sigma_1\left([w_0 \ w_1] \cdot \sigma_0\left(\begin{bmatrix} 5 & 1 & 1 \\ 5 & -1 & -1 \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}\right)\right)$$

sigmoid *e.g. sigmoid*

How to find the parameters \mathbf{w} ?

Just train the model by minimizing a corresponding loss function.

Here: binary cross-entropy since binary classification problem.



$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_{i=1}^N -\left(y_i \log f(\mathbf{x}_i, \mathbf{w}) + (1 - y_i) \log [1 - f(\mathbf{x}_i, \mathbf{w})]\right)$$

How to find the basis functions?

Different datasets require different transformations to become (almost) linearly separable.

Idea: learn the basis functions and the weights of the logistic regression **jointly** from the data (end-to-end learning)


- Previously: only learning w_{100} and w_{110}
- Now: Learn all w_{ijk} where i =layer, j = input node, k = output node

$$f(\mathbf{x}, \mathbf{W}) = \sigma_1 \left([w_{100} \ w_{110}] \cdot \sigma_0 \left(\begin{bmatrix} w_{000} & w_{010} & w_{020} \\ w_{001} & w_{011} & w_{021} \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \right) \right)$$

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \sum_{i=1}^N - \left(y_i \log f(\mathbf{x}_i, \mathbf{W}) + (1-y_i) \log [1-f(\mathbf{x}_i, \mathbf{W})] \right)$$


We get a simple Feed-Forward **Neural Network** (FFNN) with **1 hidden** layer.

Note: σ_0 and σ_1 can be arbitrary activation functions.



Making the model more complicated

Each basis function can be a more complicated function of the feature vector \mathbf{x} (a function of other basis functions rather than a function of \mathbf{x}).



By adding more hidden layers we get a 'deep' neural network:

$$f(\mathbf{x}, \mathbf{W}) = \sigma_2 \left(\mathbf{W}_2^T \sigma_1 \left(\mathbf{W}_1^T \sigma_0 \left(\mathbf{W}_0^T \mathbf{x} \right) \right) \right)$$

where $\mathbf{W} = \{\mathbf{W}_0, \mathbf{W}_1, \mathbf{W}_2\}$ are the weights to be learned.

Above architecture is called a **Multi-layered Perceptron (MLP)** = fully-connected (feed-forward) Neural Network

Why use nonlinear activation functions?

$$\underbrace{w_1 \cdot w_2 \cdot x}_{w_3 \cdot x}$$

Multiple linear layers:

$$\begin{aligned} f(\mathbf{x}, \mathbf{W}) &= \mathbf{W}_k (\mathbf{W}_{k-1} (\dots (\mathbf{W}_0 \mathbf{x}) \dots)) \\ &= (\mathbf{W}_k \mathbf{W}_{k-1} \dots \mathbf{W}_1 \mathbf{W}_0) \mathbf{x} \\ &= \mathbf{W}' \mathbf{x} \end{aligned}$$

results in a linear transformation!

Multiple nonlinear layers:

$$f(\mathbf{x}, \mathbf{W}) = \mathbf{W}_k \sigma_k(\mathbf{W}_{k-1} \sigma_{k-1}(\dots \sigma_0(\mathbf{W}_0 \mathbf{x}) \dots))$$


For non-linear functions we can not (in general) simplify it.

Activation functions

$$\sigma\left(\begin{bmatrix} 1 \\ 2 \end{bmatrix}\right) = \begin{bmatrix} \sigma(1) \\ \sigma(2) \end{bmatrix}$$


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Most activation functions are applied/operate *element-wise* when given a *multi-dimensional* input (e.g. all from above, except maxout).

Softmax activation is a notable exception!

Neural networks are universal approximators

Universal approximation theorem

An MLP with a linear output layer and one hidden layer can approximate any continuous function defined over a closed and bounded subset of \mathbb{R}^D , under mild assumptions on the activation function ('squashing' activation functions; e.g. sigmoid) and given the number of hidden units is large enough.

[Cybenko 1989; Funahashi 1989; Hornik et al 1989, 1991; Hartman et al 1990].

Also in the discrete case: a neural network can approximate any function from a discrete space to another.

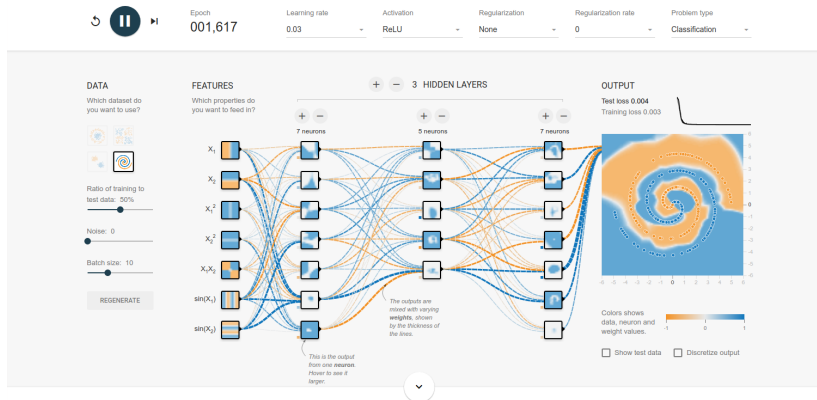
Good news: Regardless of the function we want to learn, there exists an MLP that can represent that function.

Bad news: The learning algorithm is not guaranteed to find the true parameters.

- overfitting
- picking a wrong function with 'bad' training loss

Demo: <http://playground.tensorflow.org>

Don't Worry, You Can't Break It. We Promise.



Multiple hidden layers

According to the universal approximation theorem, a two-layer feed-forward network can represent any function.

Why do we add more layers?


- The required hidden units might be exponential in the number of samples, in either discrete or continuous case.
- The issue of using the 2-layer network is easy to see in the discrete case. The number of possible functions $f : \{0, 1\}^D \rightarrow \{0, 1\}$ is 2^{2^D} . Representing all these functions requires $\mathcal{O}(2^D)$ degrees of freedom.
- For some families of functions, if we use *a few* layers we would need a large number of hidden units (and therefore parameters). But we can get the same representation power by adding *more layers*, fewer hidden units, and fewer parameters.

Multiple hidden layers

Functions that can be compactly represented with k layers may require exponentially many hidden units when using only $k - 1$ layers.

Multiple levels of latent variables allow combinatorial sharing of statistical strength.

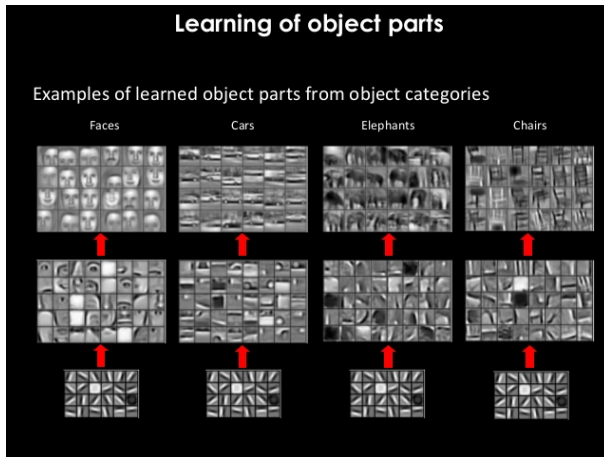
Different high-level features share lower-level features.



from: *Understanding and Improving Deep Learning Algorithms*, Yoshua Bengio, ML Google Distinguished Lecture, 2010

Multiple hidden layers

We learn “features” of “features”. This allows for *better generalization*.
(In contrast: A “wide” network tends to memorize data.)



Parameter Learning

Loss function

The choice of the loss function/cost function and the activation function of the *last* layer depend on the dataset being used and the distribution of the target variable.

Some common choices:

Output type	Output distribution	Output Layer	Cost function
Binary	Bernoulli	Sigmoid	Binary crossentropy
Discrete	Multinomial	Softmax	Crossentropy
Continuous	Gaussian	Linear	Gaussian crossentropy (Mean Squared Error)
Continuous	Arbitrary	GAN, VAE, ...	Various

Example 1: Binary classification

Given a set of labeled data $\{\mathbf{x}_i, y_i\}_{i=1}^N$, where $y_i \in \{0, 1\}$.
Usually one takes an NN with sigmoidal outputs

$$y = \sigma(a) = \frac{1}{1 + \exp(-a)} = f(\mathbf{x}, \mathbf{W})$$

In this case, we would use the binary cross-entropy between $f(\mathbf{x}_i, \mathbf{W})$ and y_i as the cost function

$$E(\mathbf{W}) = - \sum_{i=1}^N \left(y_i \log f(\mathbf{x}_i, \mathbf{W}) + (1 - y_i) \log [1 - f(\mathbf{x}_i, \mathbf{W})] \right)$$

$$\min_{\mathbf{W}} E(\mathbf{W})$$

Example 2: Standard multi-class classification

Take a set of labeled data $\{\mathbf{x}_i, y_i\}_{i=1}^N$, where $y_i \in \{0, 1\}^K$ (i.e. 1-of-K coding).

For this, we usually take an NN with Softmax outputs

$$y_k = \frac{\exp(a_k)}{\sum_j \exp(a_j)} = f_k(\mathbf{x}, \mathbf{W})$$

In this case, we would use cross-entropy between $f(\mathbf{x}_i, \mathbf{W})$ and y_i as the cost function

$$E(\mathbf{W}) = - \sum_{n=1}^N \sum_{k=1}^K \left(y_{nk} \log f_k(\mathbf{x}_n, \mathbf{W}) \right)$$

Minimizing the cost function


In practice $E(\mathbf{W})$ is often non-convex \implies optimization is tricky

- a local minimum is not necessarily a global minimum.
- potentially there exist several local minima, many of which can be equivalent (see next tutorial session)
- often it is not possible to find a global minimum nor is it useful.

We may find a few local minima, and pick the one with higher performance on a validation set.

Default approach:
find a local minimum by
using gradient descent

$$\mathbf{W}^{(t+1)} = \mathbf{W}^{(t)} - \alpha \nabla_{\mathbf{W}} E(\mathbf{W}^{(t)})$$




How can we compute the gradient?

- 1) By hand: manually working out $\nabla_{\mathbf{W}} E$ and coding it is tricky and cumbersome (furthermore: see point 3).
- 2) Numeric: Can be done as

$$\frac{\partial E_n}{\partial w_{ij}} = \frac{E_n(w_{ij} + \epsilon) - E_n(w_{ij})}{\epsilon} + \mathcal{O}(\epsilon)$$

- Each evaluation of the above equation roughly requires $\mathcal{O}(|W|)$ operations, where $|W|$ is the dimensionality of weight space.
- The evaluation has to be done for each parameter independently. Therefore computing $\nabla_{\mathbf{W}} E$ requires $\mathcal{O}(|W|^2)$ operations!

How can we compute the gradient?



3) Symbolic differentiation: Automates essentially how you would compute the gradient function by hand.

But: Explicitly 'writing down' (and computing) the gradient function for every parameter is very expensive

- potentially exponentially many different cases (e.g. when having multiple layers with ReLUs)
- many terms reappear in the gradient computation for *different* parameters (since the function f is hierarchically constructed); these terms could be re-used to make computation faster; however symbolic differentiation does not exploit this insight

$$\frac{\delta f}{\delta w_{i,j,k}}(x) = \left\{ \begin{array}{l} \vdots \\ \vdots \end{array} \right.$$

4) Automatic differentiation: e.g. **backpropagation** for neural networks

- Computes $(\nabla_{\mathbf{W}} E)$ automatically and efficiently
- Evaluation in $\mathcal{O}(|W|)$ (every 'neuron' is visited only twice)

Recap: Chain Rule

Recall: Chain rule of calculus

- for scalars $x, y, z \in \mathbb{R}$

$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$$

- for vectors $\mathbf{x} \in \mathbb{R}^m, \mathbf{y} \in \mathbb{R}^n, z \in \mathbb{R}$

$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i}$$

or with a more compact notation (using Jacobian matrix):

$$\nabla_{\mathbf{x}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^T \nabla_{\mathbf{y}} z$$

$$\text{where } \left[\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right]_{ij} = \frac{\partial y_i}{\partial x_j}$$

Helpful Concept: Computational Graph

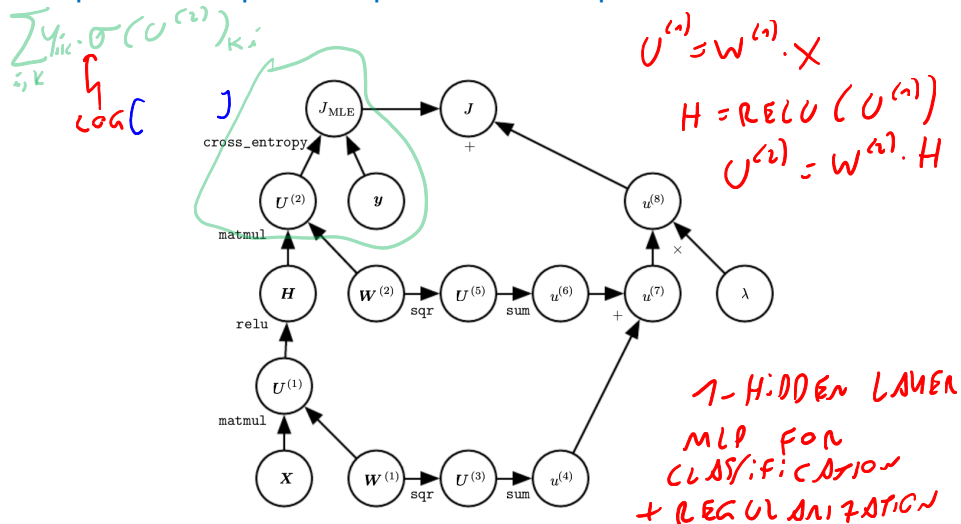



Figure 6.11: The computational graph used to compute the cost used to train our example of a single-layer MLP using the cross-entropy loss and weight decay.

Backpropagation: Toy Example


Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?

$$d = x + y$$
$$f = d \cdot z$$



Backpropagation: Toy Example

Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?



$$d = x + y$$

$$f = d \cdot z$$


$$\frac{\delta f}{\delta z} = d = -2$$

$$\frac{\delta f}{\delta d} = z = 4$$

$$\frac{\delta f}{\delta x} = \frac{\delta f}{\delta d} \cdot \frac{\delta d}{\delta x} = 4 \cdot 1 = 4$$

Backpropagation: Toy Example


Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?



$$\frac{\delta f}{\delta f} = 1$$


Backpropagation: Toy Example

Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?




Backpropagation: Toy Example

Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?



Backpropagation: Toy Example

Example: $f = (x + y) \cdot z$. Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$?



Backpropagation

In practice, for most of the cost functions we have that

$$E(\mathbf{W}) = \sum_{n=1}^N E_n(\mathbf{W})$$

The specific structure of the feed-forward neural networks allows us to efficiently compute $\nabla_{\mathbf{W}} E_n(\mathbf{W})$ for each instance in the training set.

For simplicity, let's consider only the element-wise activation functions (let's ignore softmax for now).


Notation:

- $w_{\ell ij}$: a weight where ℓ =layer, i = input node and j = output node
- $z_{\ell i}^{(n)}$: the value of a neuron, ℓ =layer, i = node index, n = instance
- $a_{\ell i}^{(n)}$: the value of a logit, ℓ =layer and i = node index, n = instance; $\mathbf{a}_{\ell}^{(n)} = \mathbf{W}_{\ell-1} \cdot \mathbf{z}_{\ell-1}^{(n)}$
- $h_{\ell}(\cdot)$: the activation function of the ℓ -th layer; $z_{\ell j}^{(n)} = h_{\ell}(a_{\ell j}^{(n)})$

Backpropagation

$$\frac{\partial E_n}{\partial w_{(\ell-1)ij}} = \boxed{\frac{\partial E_n}{\partial a_{\ell j}^{(n)}}} \cdot \boxed{\frac{\partial a_{\ell j}^{(n)}}{\partial w_{(\ell-1)ij}}}$$

δ z



$$a = z \cdot w$$

$$\frac{\delta a}{\delta w} = z$$

- $\delta_{\ell j}^{(n)} \equiv \frac{\partial E_n}{\partial a_{\ell j}^{(n)}}$, is called the error of the j -th neuron at the ℓ -th layer
- The term $\frac{\partial a_{\ell j}^{(n)}}{\partial w_{(\ell-1)ij}}$ is simply equal to $z_{(\ell-1)i}^{(n)}$

Algorithm 1: Backpropagation

- 1 For each instance \mathbf{x}_n in the training set:
- 2 Forward pass: compute the values of $\mathbf{z}_1^{(n)}, \mathbf{z}_2^{(n)}, \dots, \mathbf{z}_L^{(n)}$ (and $\mathbf{a}_l^{(n)}$)
- 3 Compute the errors recursively as $\delta_L^{(n)}, \delta_{(L-1)}^{(n)}, \dots, \delta_2^{(n)}$
- 4 Compute $\nabla_{\mathbf{W}} E_n$ using the above equation.
- 5 Compute the gradient as $\nabla_{\mathbf{W}} E = \sum_n (\nabla_{\mathbf{W}} E_n)$

Forward pass

The forward pass is trivial. Just evaluate the function.

$$\mathbf{z}_0^{(n)} = \mathbf{x}^{(n)}$$

$$\mathbf{a}_\ell^{(n)} = \mathbf{W}_{\ell-1} \mathbf{z}_{\ell-1}^{(n)}$$

$$\mathbf{z}_\ell^{(n)} = \mathbf{h}_\ell(\mathbf{a}_\ell^{(n)})$$

Backward pass

$$z = h(a)$$

Computing the error recursively:

$$\frac{\partial E_n}{\partial a_{\ell j}^{(n)}} = \frac{\partial E_n}{\partial a_{(\ell+1)j}^{(n)}} \cdot \frac{\partial a_{(\ell+1)j}^{(n)}}{\partial z_{\ell j}^{(n)}} \cdot \frac{\partial z_{\ell j}^{(n)}}{\partial a_{\ell j}^{(n)}} \Rightarrow \delta_{\ell j}^{(n)} = (\delta_{\ell+1}^{(n)})^T \mathbf{W}_{\ell j} : h'_{\ell}(a_{\ell j}^{(n)})$$

Derivative of the activation function is usually known.


For example, for the sigmoid function: $\sigma'(a) = \sigma(a)(1 - \sigma(a))$

Computing δ for the last layer

To make it simple, let's assume the following:

1. the last layer is linear, i.e., $h_L(a) = a$
2. we've used the mean squared error (we are solving, e.g., a regression problem)

$$\begin{aligned}\delta_{Lj}^{(n)} &= \frac{\partial E_n}{\partial a_{Lj}^{(n)}} = \frac{\partial E_n}{\partial z_{Lj}^{(n)}} \frac{\partial z_{Lj}^{(n)}}{\partial a_{Lj}^{(n)}} \\ &= \frac{\partial}{\partial z_{Lj}^{(n)}} \left(\frac{1}{2} (z_{Lj}^{(n)} - y_j^{(n)})^2 \right) \times 1 \\ &= z_{Lj}^{(n)} - y_j^{(n)} \Rightarrow \text{error!}\end{aligned}$$



We can follow a similar procedure in the case of logistic sigmoid activation function along with binary cross entropy, or softmax activation function along with cross entropy.

Backpropagation: Summary

$$\frac{\delta f}{\delta w_{i,j,k}}$$

- Allows to compute the gradient very efficiently
- Only two passes through the computational graph required (forward, backward)
- We only need to know the 'local' gradient function per node/activation function in the computational graph
 - above, we assumed to know $h'_{\ell}(a_{\ell j}^{(n)})$
 - as well as the gradient of the loss/cost function
- If you use any state-of-the-art deep learning framework, for every operation/activation/loss function, the corresponding 'local' gradient is implemented as well
 - given this, you can construct arbitrary complex functions and the gradient can be computed automatically

Overall, we can now easily apply our well-known gradient descent to learn the optimal parameters.


Will we reach a local minimum?

Cliffs: In some points, the norm of the gradient might become too large.

- The next step will catapult the parameters very far.

To alleviate this issue, there is a simple yet effective trick: Gradient clipping

- We can clip the elements of the gradient to an interval $[-c, +c]$




source: Goodfellow

Will we reach a local minimum?



The cost function might be constant over wide flat regions of the weight space.

- Gradient descent actually fails, because it cannot find the steepest direction.
- Prevalent when using sigmoid activation function: its input values might be too large or too small \Rightarrow zero gradient for incoming weights (**vanishing gradient**)



ReLU activation function alleviates the above problem, because the gradient is always 1 at least on positive numbers.

Dead ReLU units: Assume that the input of a ReLU unit becomes negative for all data instances because of, e.g., a large negative bias. The gradient w.r.t. the incoming weights becomes zero \Rightarrow The unit will remain at this state forever.



Will we reach a local minimum? ... $(w \cdot (\text{ReLU}(w \cdot x)))$

Repetition of a parameter might also result in vanishing gradient problem (prevalent in RNNs; see later).

Consider the following simple case:

- The operation is multiplying the input vector by the matrix \mathbf{W} , t times

- To put it simple:

$$\mathbf{W} = \mathbf{V} \text{diag}(\mathbf{D}) \mathbf{V}^{-1} \Rightarrow \mathbf{W}^t = \mathbf{V} (\text{diag}(\mathbf{D}))^t \mathbf{V}^{-1}.$$

$w^{t \cdot x}$

In this simple case, the gradient w.r.t. the elements of \mathbf{D} most probably would either vanish or explode:

- if $\mathbf{D}_{ii} < 1.0$, then \mathbf{D}_{ii}^t would be near zero, the effect of the gradient would be faded.
- if $\mathbf{D}_{ii} > 1.0$, then \mathbf{D}_{ii}^t would explode and makes the computations unstable.